

Cloud



Caltech | Center for Technology & Management Education

Post Graduate Program in Cloud

Cloud



Caltech

**Center for Technology &
Management Education**

AWS Solution Architect: Associate Level



Amazon Elastic Compute Cloud

Learning Objectives

By the end of the lesson, you will be able to:


- 🕒 Explain Amazon EC2 and its uses
- 🕒 Describe Amazon Machine Image (AMI)
- 🕒 Explain the difference between EBS and EFS
- 🕒 Describe the use of EBS snapshots
- 🕒 Describe the different instance pricing options
- 🕒 Illustrate Auto Scaling, FSx, and HPC
- 🕒 Evaluate the AWS recommended best practices for EC2



Introduction to Amazon EC2

Elastic Compute Cloud (EC2)

Amazon Elastic Compute Cloud (Amazon EC2) is a web service that provides scalable computing capacity in the Amazon Web Services (AWS) cloud.



Benefits of Using EC2:

Elastic Web-Scale Computing


Flexible Cloud Hosting
Services

AWS Integration

Reliability and Security

Low Cost

- EC2 increases or decreases your storage capacity in minutes.
- It launches thousands of server instances simultaneously.



Benefits of Using EC2:

Elastic Web-Scale Computing


Flexible Cloud Hosting
Services

AWS Integration

Reliability and Security

Low Cost

EC2 launches numerous operating systems, instance types, and software in just minutes.



Benefits of Using EC2:

Elastic Web-Scale Computing


Flexible Cloud Hosting
Services

AWS Integration

Reliability and Security

Low Cost

EC2 is integrated with other AWS products such as Amazon S3, Amazon RDS, and Amazon SQS, to provide a complete IT architecture solution.



Benefits of Using EC2:

Elastic Web-Scale Computing

Flexible Cloud Hosting
Services

AWS Integration

Reliability and Security

Low Cost

- AWS operates a Service Level Agreement (SLA) commitment of 99.95% availability.
- With Amazon VPC, you can easily create secure and robust networks to run your Amazon EC2 instances.



Benefits of Using EC2:

Elastic Web-Scale Computing

Flexible Cloud Hosting
Services

AWS Integration

Reliability and Security

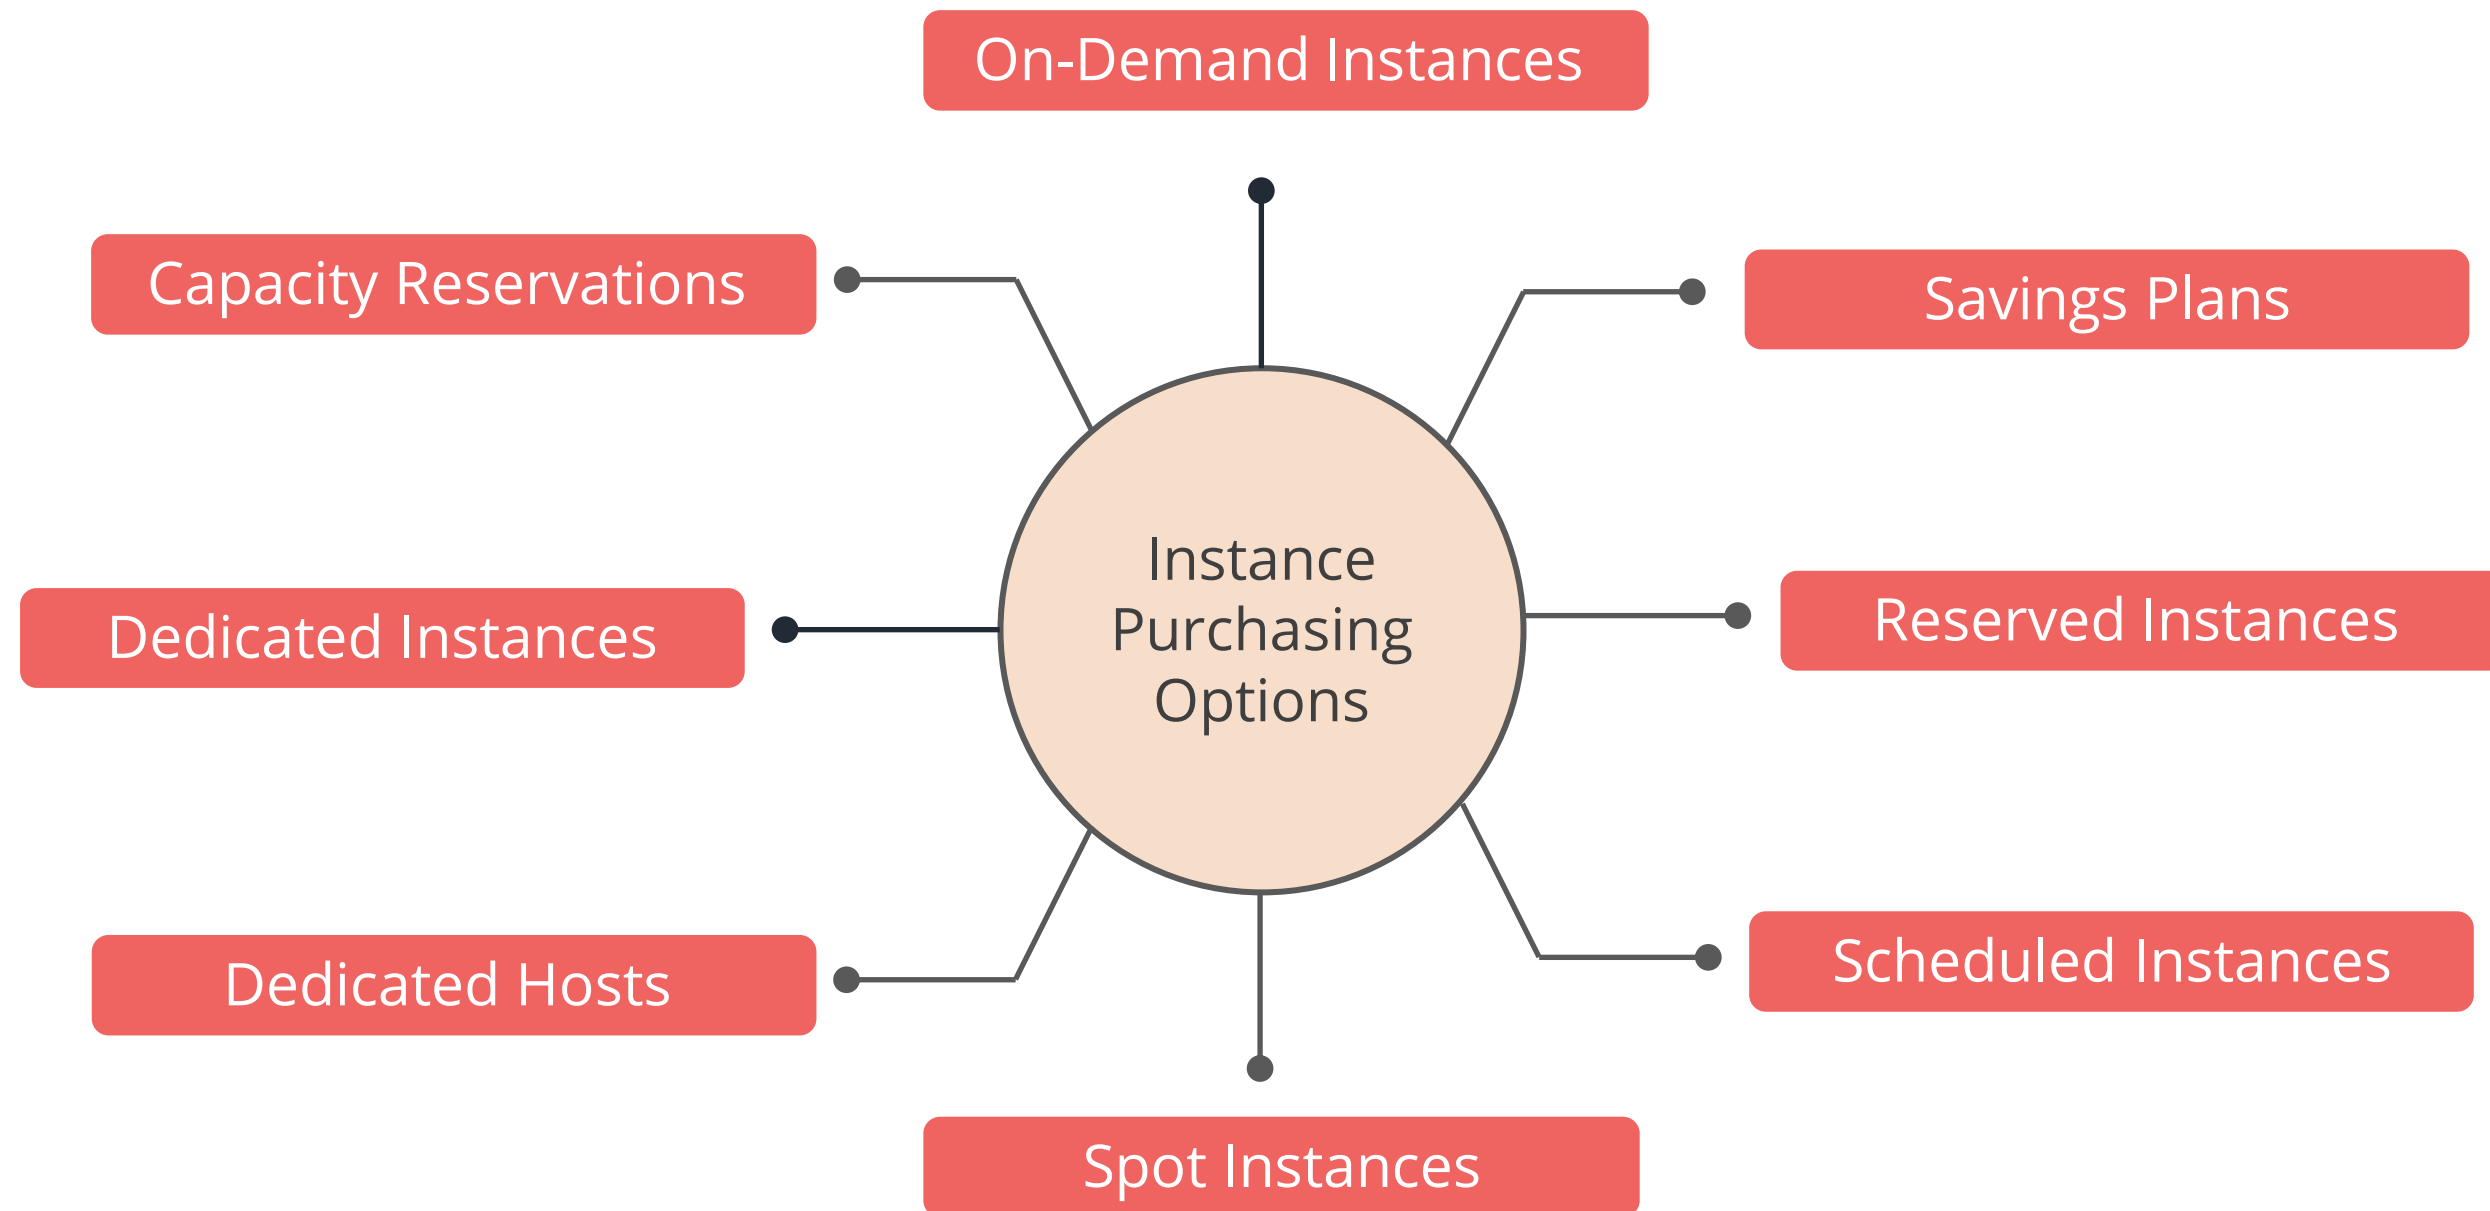
Low Cost

- AWS charges you by seconds and you only pay for what you use.
- Rates are lower than your existing on-premise infrastructure.



EC2 Instance Purchasing Options

Amazon EC2 provides the following purchasing options to enable you to optimize your costs based on your needs:



Amazon Machine Images (AMI)

Amazon Machine Image (AMI)

Amazon's definition of an AMI:

"An Amazon Machine Image (AMI) provides the information required to launch an instance. You specify an AMI when you launch an instance, and you can launch as many instances from as many different AMIs as you need."



Amazon Machine
Image (AMI)

Amazon Machine Image (AMI)

AMI is a virtual instance that includes:

- A template for the root volume for the instance
- Launch permissions to control AMI launch instances
- A block device mapping that specifies volumes to attach to the instance



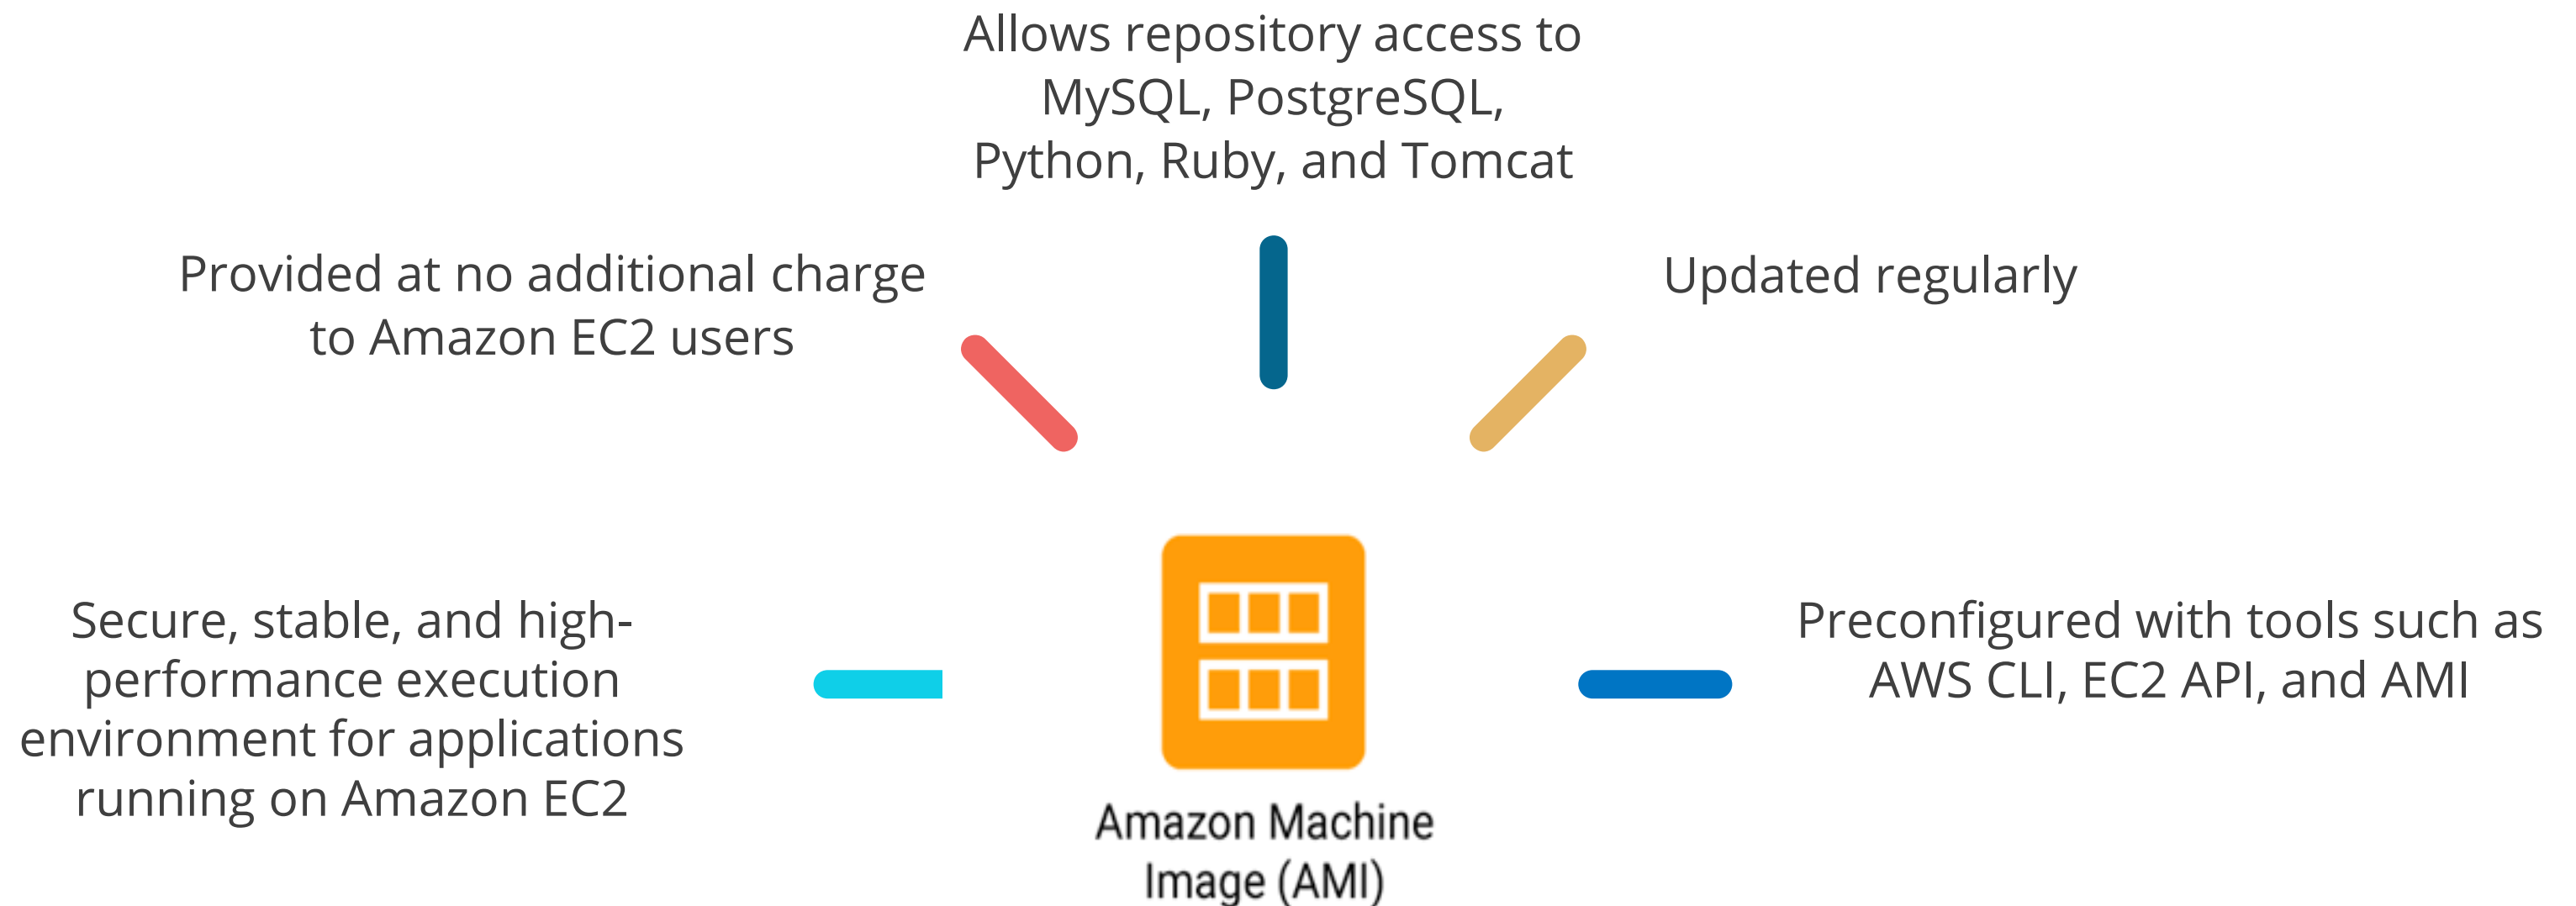
AMI



Amazon Linux AMI


Amazon Linux AMI is a supported and maintained Linux image provided by AWS.

The following are the features of the Amazon Linux AMI:



Choose an AMI

The first step of launching any new instance is selecting an AMI.



Services ▾Resource Groups ▾📌

🔔Asmita-RayAsmita-Ray @ 3513... ▾N. Virginia ▾Support ▾

1. Choose AMI2. Choose Instance Type3. Configure Instance4. Add Storage5. Add Tags6. Configure Security Group7. Review

Step 1: Choose an Amazon Machine Image (AMI)

[Cancel and Exit](#)

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. You can select an AMI provided by AWS, our user community, or the AWS Marketplace; or you can select one of your own AMIs.

[Search by Systems Manager parameter](#)


Quick Start

My AMIs

AWS Marketplace

Community AMIs

☐ Free tier only ⓘ




Amazon Linux

Free tier eligible

Amazon Linux 2 AMI (HVM), SSD Volume Type - ami-08f3d892de259504d (64-bit x86) / ami-0ba960472fc891755 (64-bit Arm)

64-bit (x86)64-bit (Arm)

Amazon Linux 2 comes with five years support. It provides Linux kernel 4.14 tuned for optimal performance on Amazon EC2, systemd 219, GCC 7.3, Glibc 2.26, Binutils 2.29.1, and the latest software packages through extras.
Root device type: ebsVirtualization type: hvmENA Enabled: Yes



Amazon Linux

Free tier eligible

Amazon Linux AMI 2018.03.0 (HVM), SSD Volume Type - ami-0e9089763828757e1

64-bit (x86)

The Amazon Linux AMI is an EBS-backed, AWS-supported image. The default image includes AWS command line tools, Python, Ruby, Perl, and Java. The repositories include Docker, PHP, MySQL, PostgreSQL, and other packages.
Root device type: ebsVirtualization type: hvmENA Enabled: Yes

 Feedback

 English (US)

© 2008 - 2020, Amazon Web Services, Inc. or its affiliates. All rights reserved. [Privacy Policy](#) [Terms of Use](#)

Caltech

Center for Technology & Management Education

simplilearn

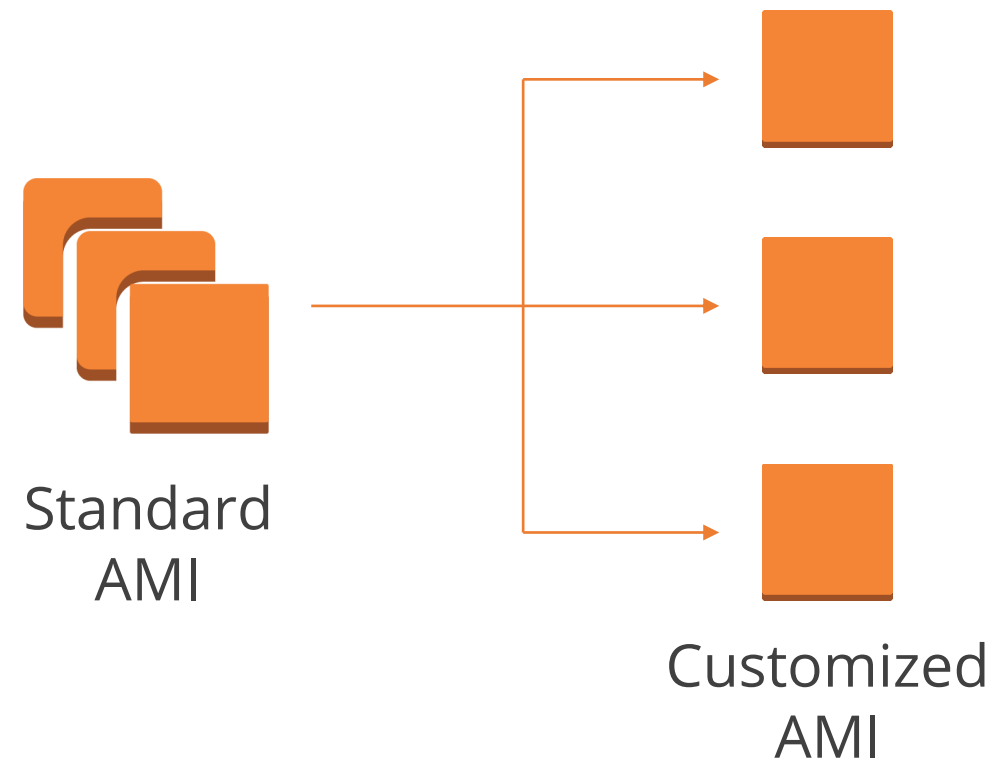
AWS Marketplace

The AWS Marketplace is an online store where customers can find, buy, deploy, and manage the software, data, and services that run on AWS from third-party vendors.

The screenshot displays the AWS Marketplace interface. At the top, the AWS logo is on the left, and navigation links for Services, Resource Groups, and a user profile (Asmita-Ray) are on the right. Below the header, the left sidebar contains links for Manage subscriptions, Discover products, and Product Support Connection. The main content area is titled 'AWS Marketplace > Discover products > Search results'. It features a 'Refine results' section on the left with filters for Categories (Infrastructure Software, DevOps, Business Applications, Machine Learning, Data Products, IoT, Industries) and Pricing plans (Annual, Hourly, Free, Monthly, Bring Your Own License, By Units, By Users, By Requests). The search results section on the right shows 'Search AWS Marketplace products' with a search bar and a list of products. The first product is 'Microsoft Windows Server 2019 Base' by Amazon Web Services, version 2020.06.10, with a 5-star rating and 2 reviews. The description states: 'Amazon EC2 running Microsoft Windows Server is a fast and dependable environment for deploying applications using the Microsoft Web Platform. Amazon EC2 enables you to run compatible Windows-based solutions on AWS' high-performance, reliable, cost-effective, cloud computing platform.'

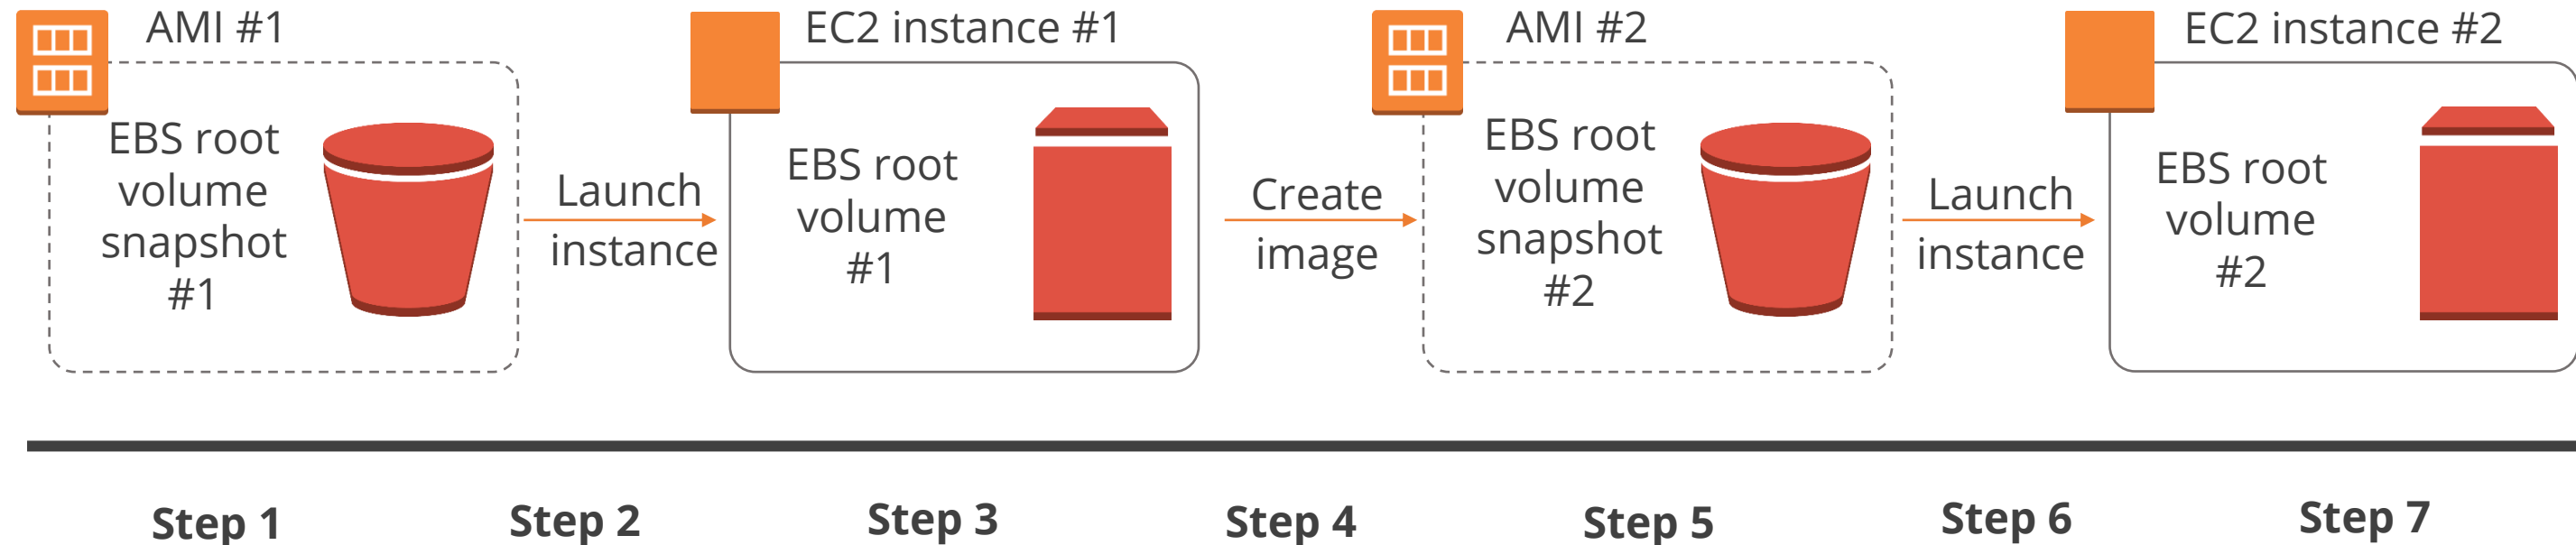
My AMIs

Customize the instance that you launch from a public AMI and then save that configuration as a custom AMI for your use.



Creating AMIs


Use an existing AMI, launch an instance, customize it, create a new AMI from it, and finally launch an instance of your new AMI.



AMI Distribution

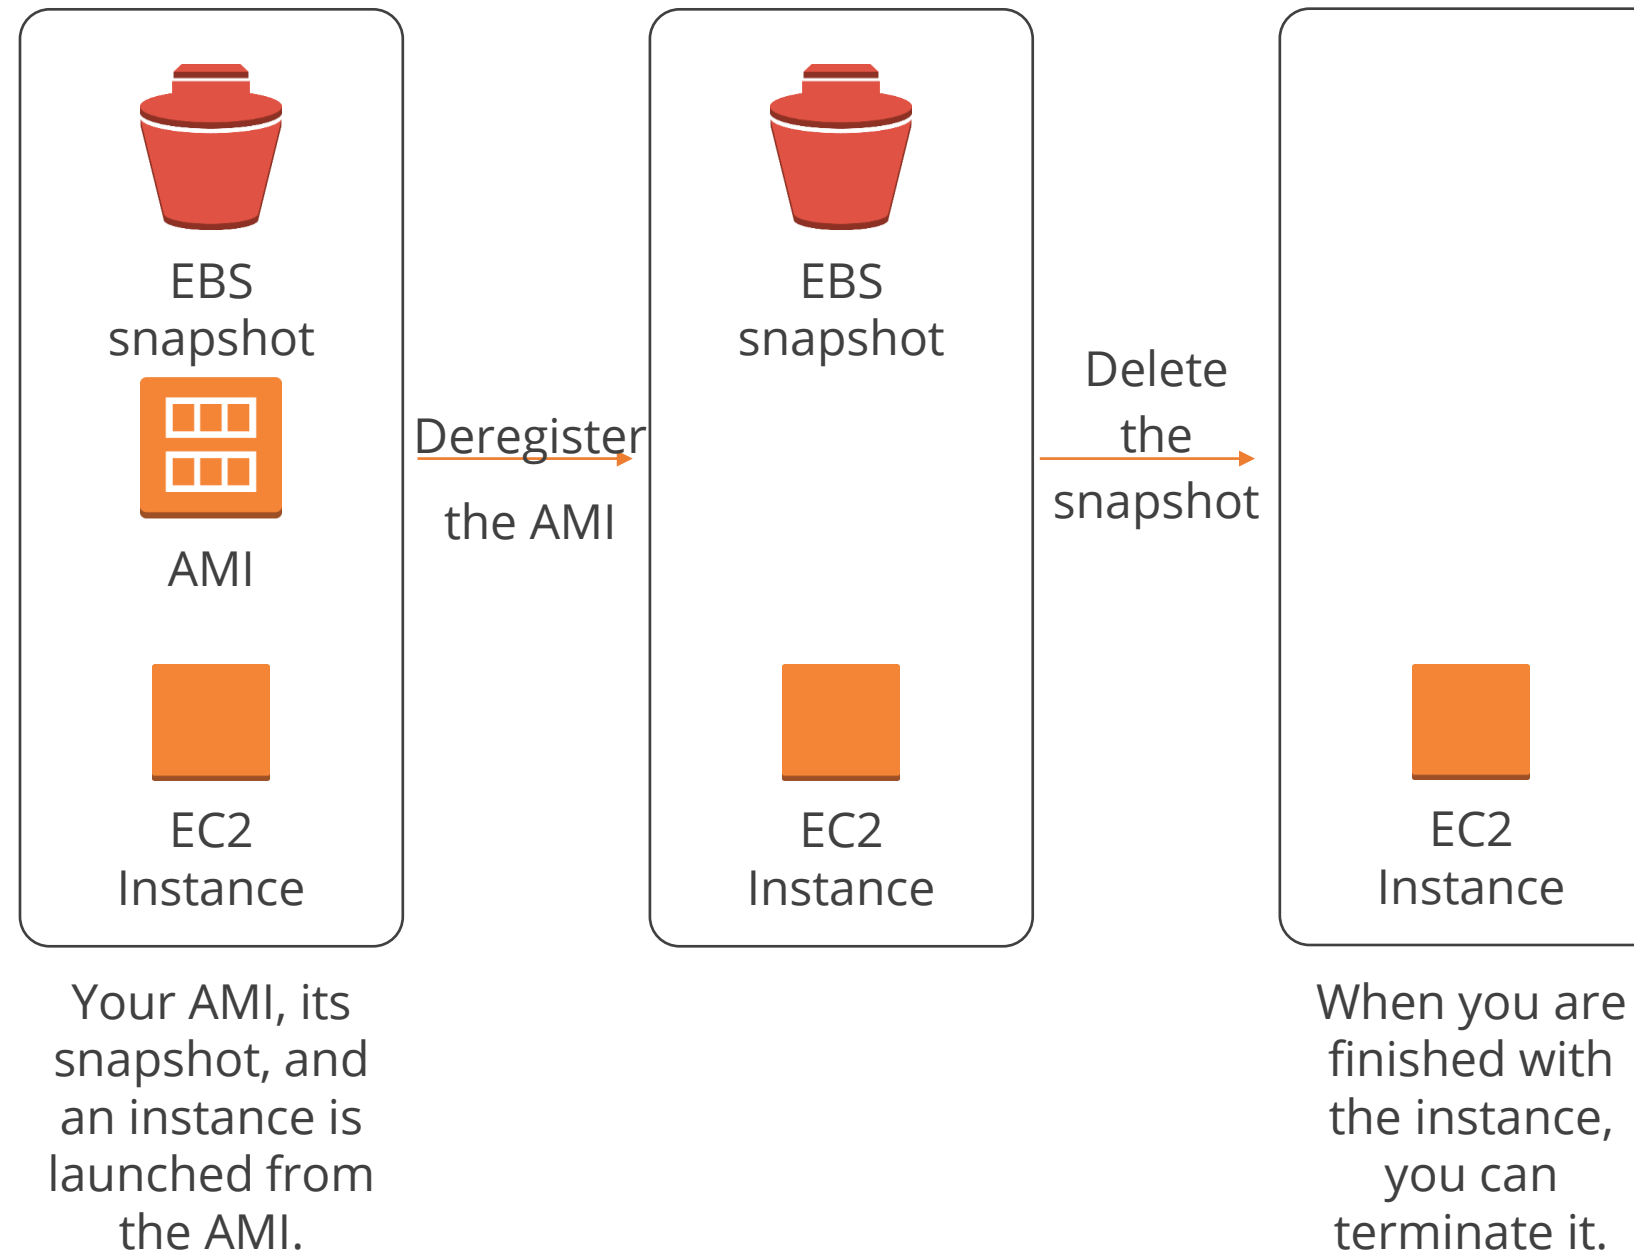
An AMI can be:

- Kept private
- Shared with a specific list of AWS accounts
- Made public



Deregistering AMIs

An AMI can be deregistered after its work is done. Once deregistered, it cannot launch new instances.



Assisted Practice

Assisted Practice: Create a Linux-based EC2 Instance

Duration: 30 min.

Problem Statement:

You are given a project to create a Linux-based EC2 instance.

Assisted Practice: Guidelines to Create a Linux-based EC2 Instance

Steps to perform:

1. Open the Amazon EC2 console
2. Create an EC2 Instance and launch the instance
3. Convert your private key to PuTTY using PuTTYgen
4. Connect to your Linux Instance

Assisted Practice

Create a Windows-based EC2 Instance

Duration: 20 min.

Problem Statement:

You are given a project to create a Windows-based EC2 instance.

Assisted Practice: Guidelines to Create a Windows-based EC2 Instance

Steps to perform:

1. Open the Amazon EC2 console
2. Create an EC2 Instance and launch the instance
3. Click on Download Remote Desktop File and save it
4. Connect to your Windows Instance

Assisted Practice

Create a Custom AMI

Duration: 10 min.

Problem Statement:

You are given a project to create a custom AMI.

Assisted Practice: Guidelines to Create a Custom AMI

Steps to perform:

1. Create image commands from the Actions dropdown menu
2. Give the detailed information of the image and create it
3. Check the custom image that you have created

Assisted Practice

Create an Instance Using a Custom AMI

Duration: 25 min.

Problem Statement:

You are given a project to create an instance using a custom AMI.

Assisted Practice: Guidelines to Create a Custom AMI

Steps to perform:

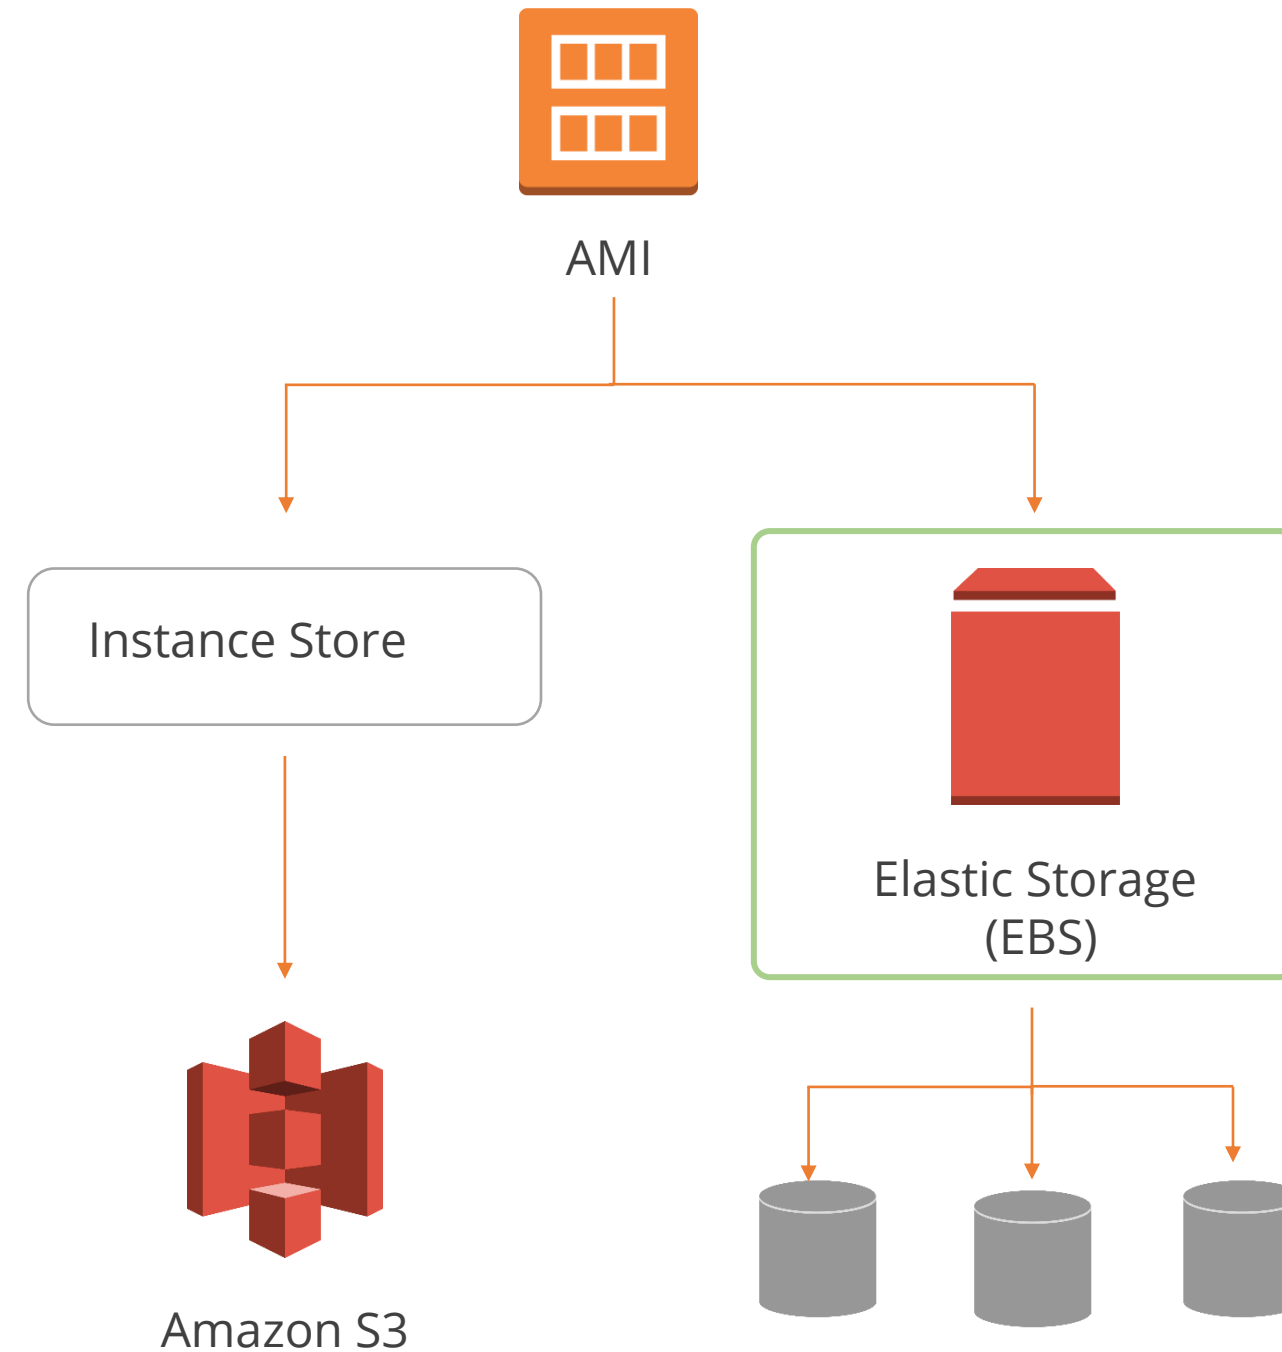
1. Open the Amazon EC2 console
2. Choose AMIs from the navigation bar
3. Launch the instance using a custom AMI

Introduction to Amazon EBS

Root Device Storage

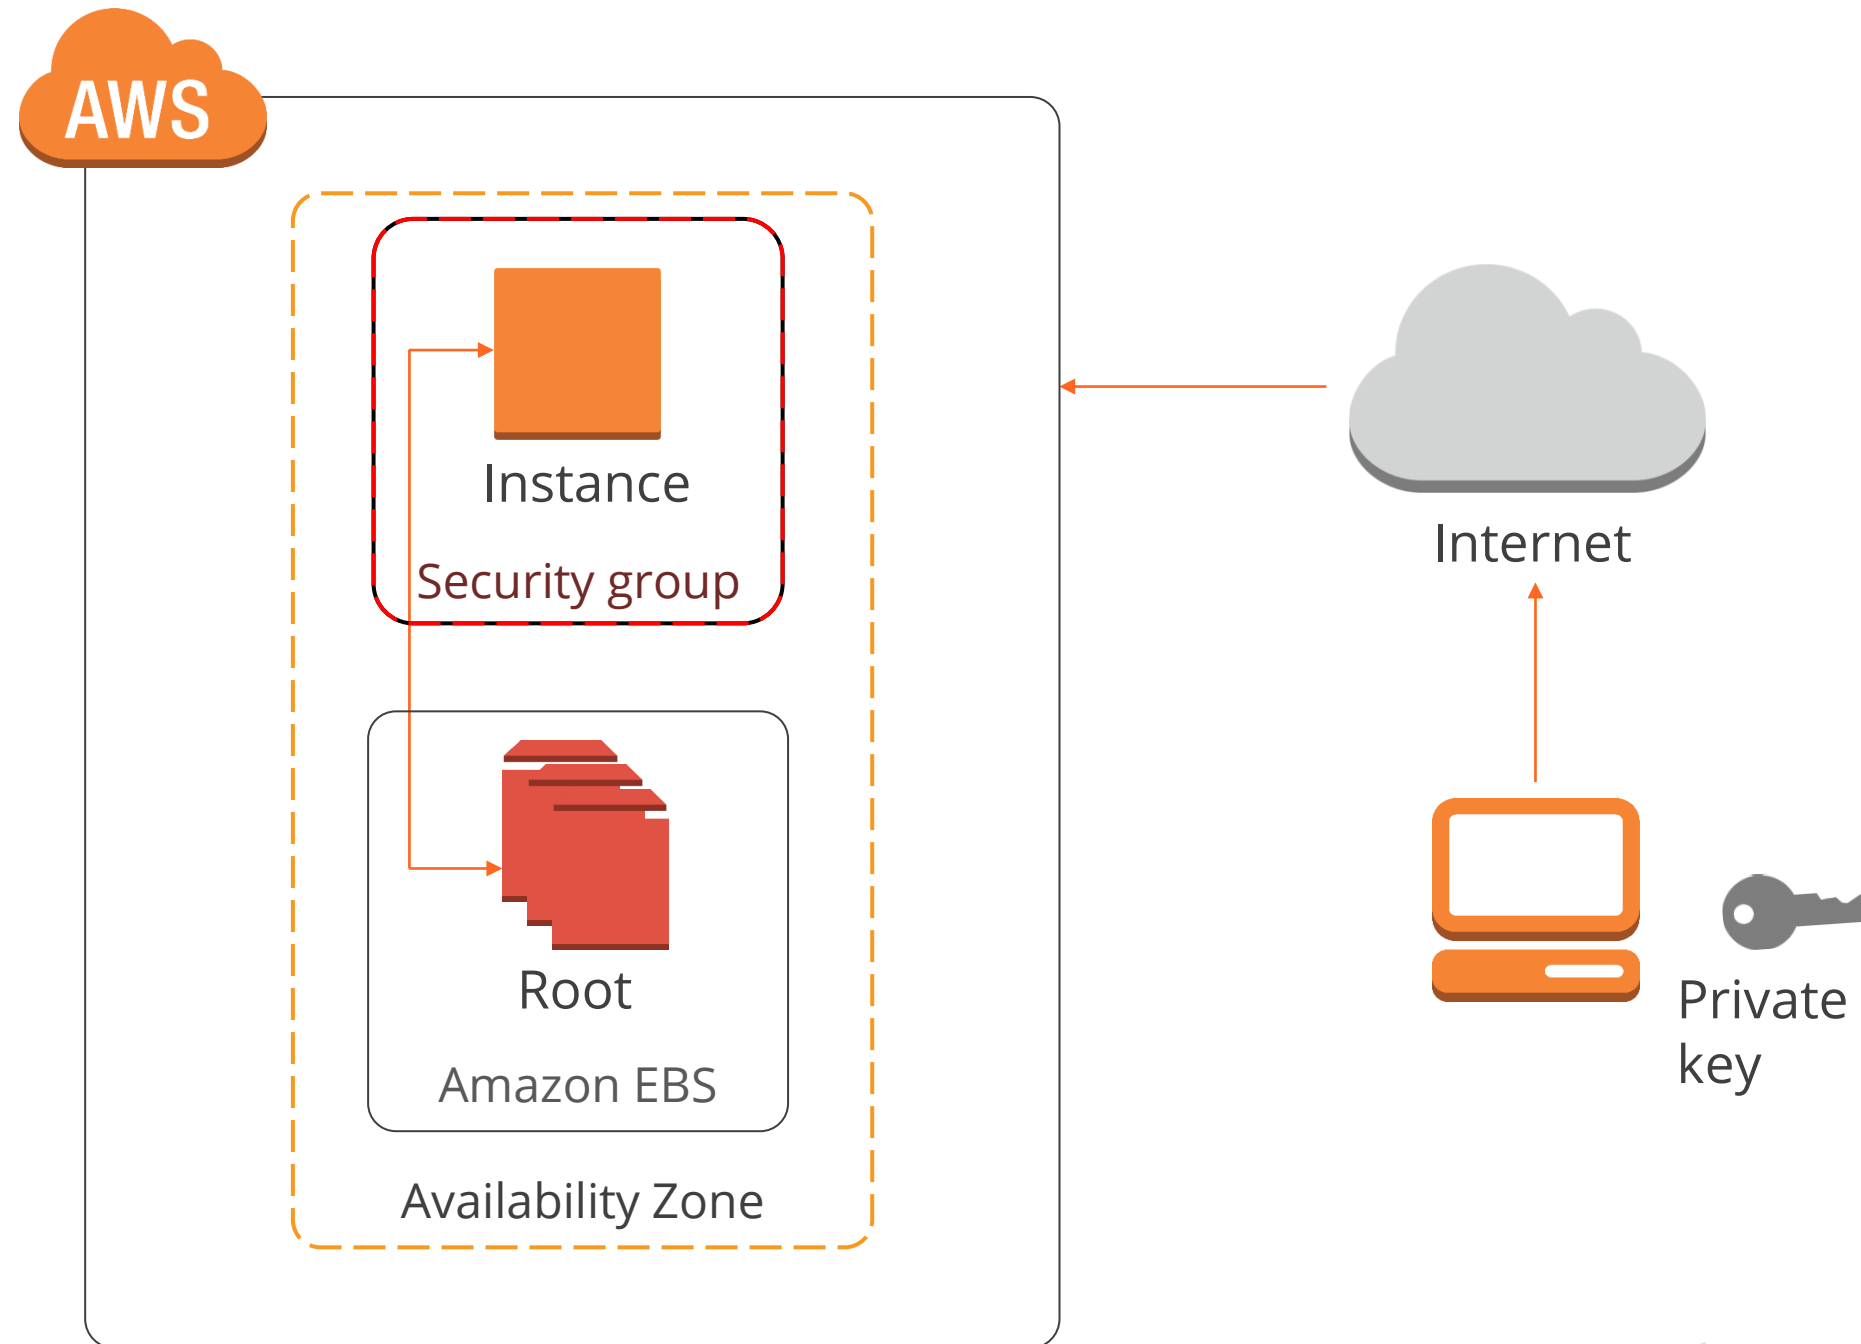
There are two types of root device storage for AMIs:

- Instance store
- Amazon EBS

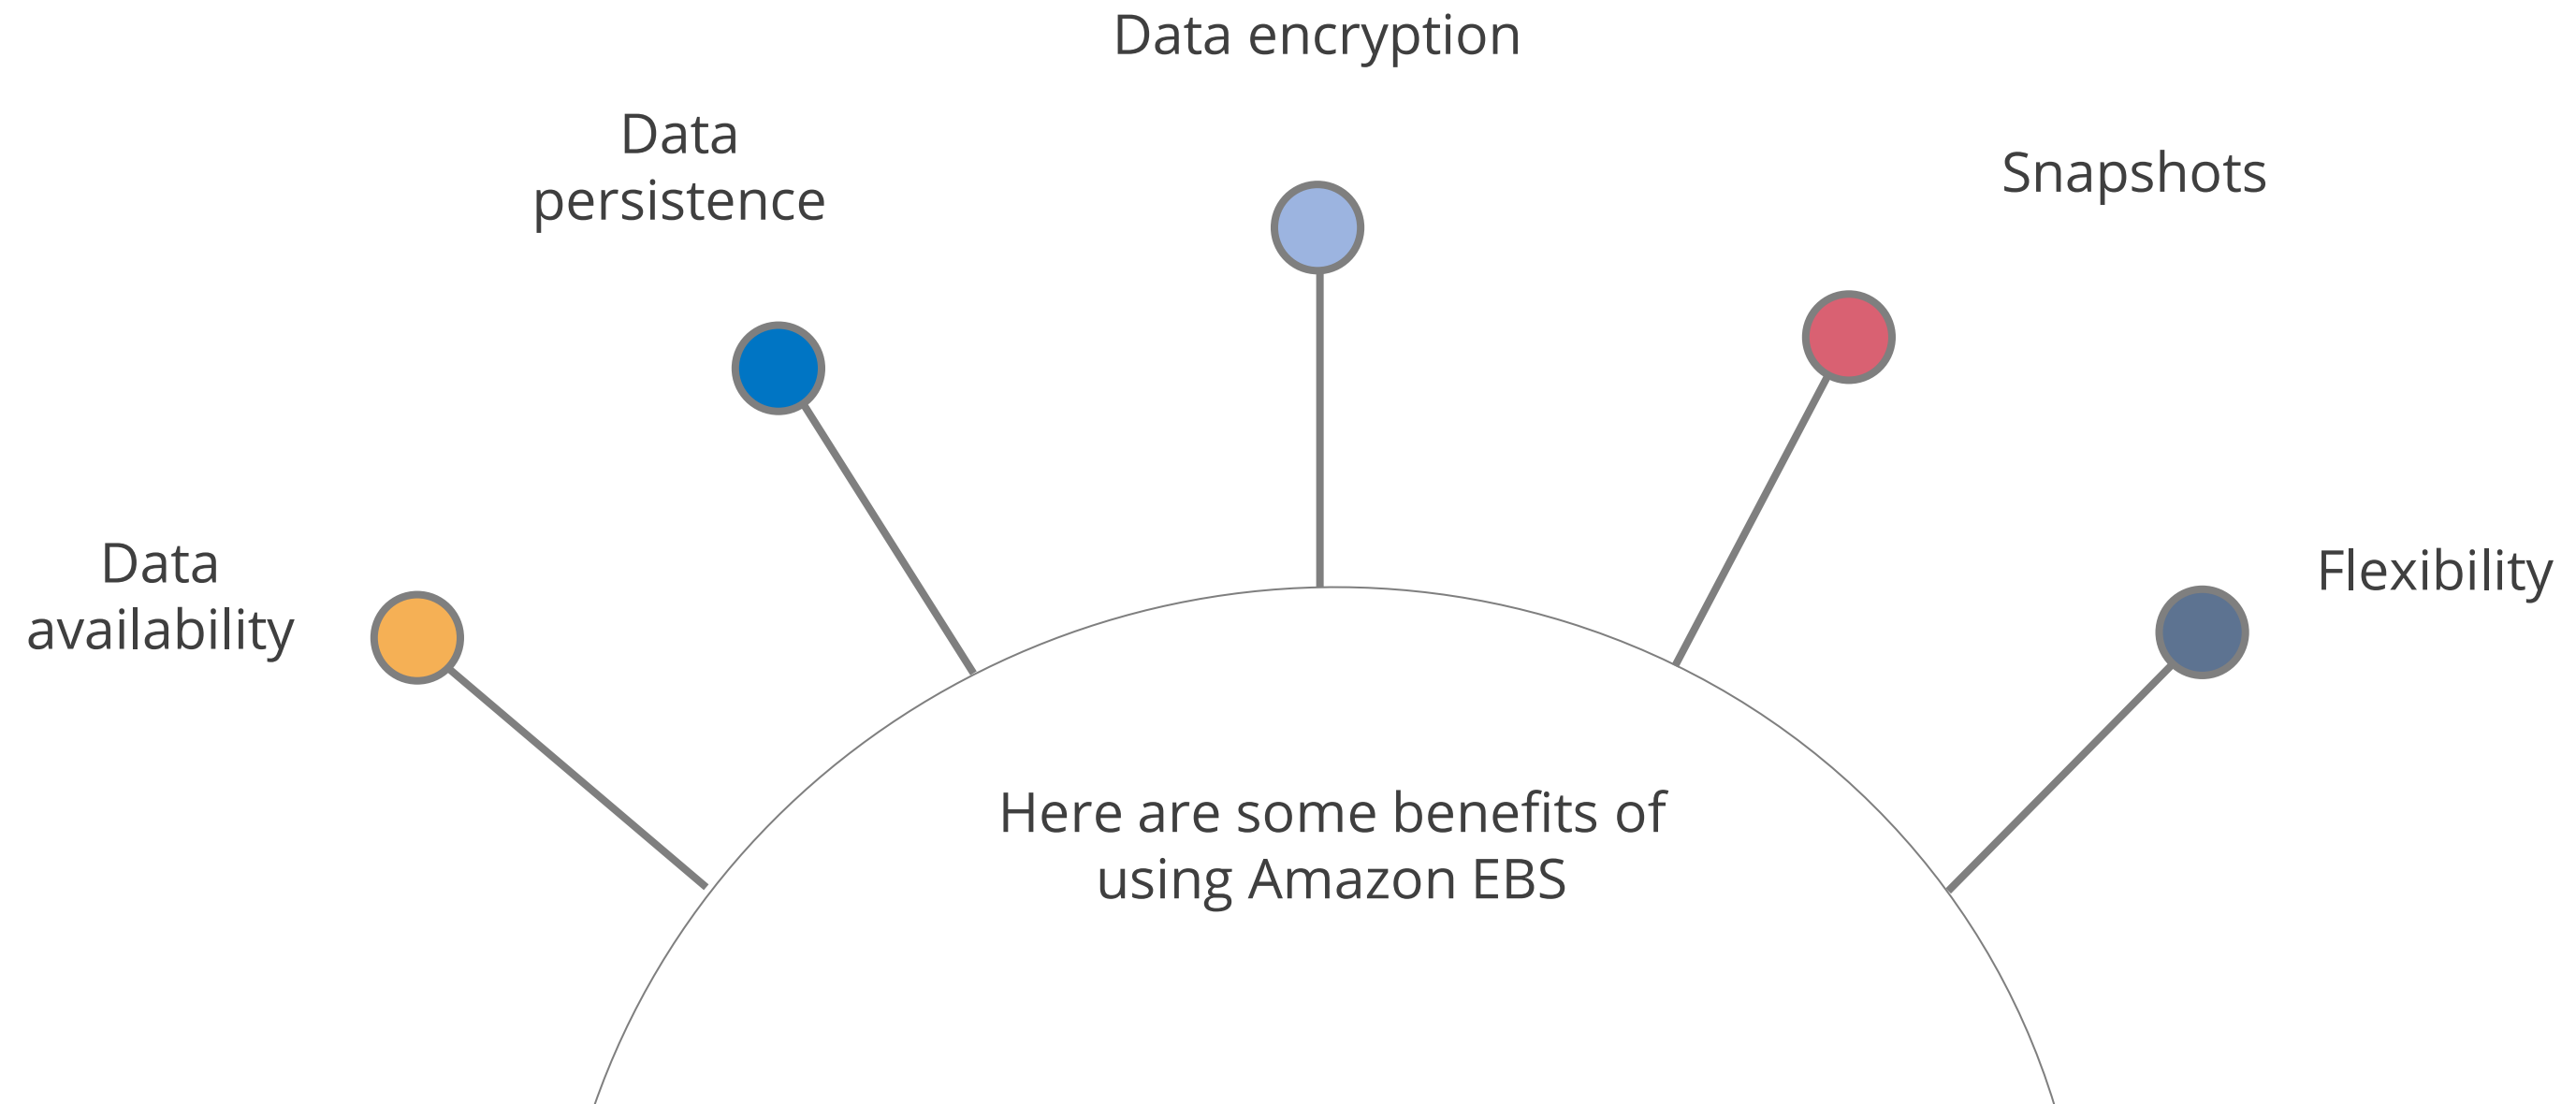


EBS Overview

Amazon Elastic Block Store (Amazon EBS) provides persistent block-level storage volumes for use with Amazon EC2 instances in the AWS cloud.



Benefits of EBS

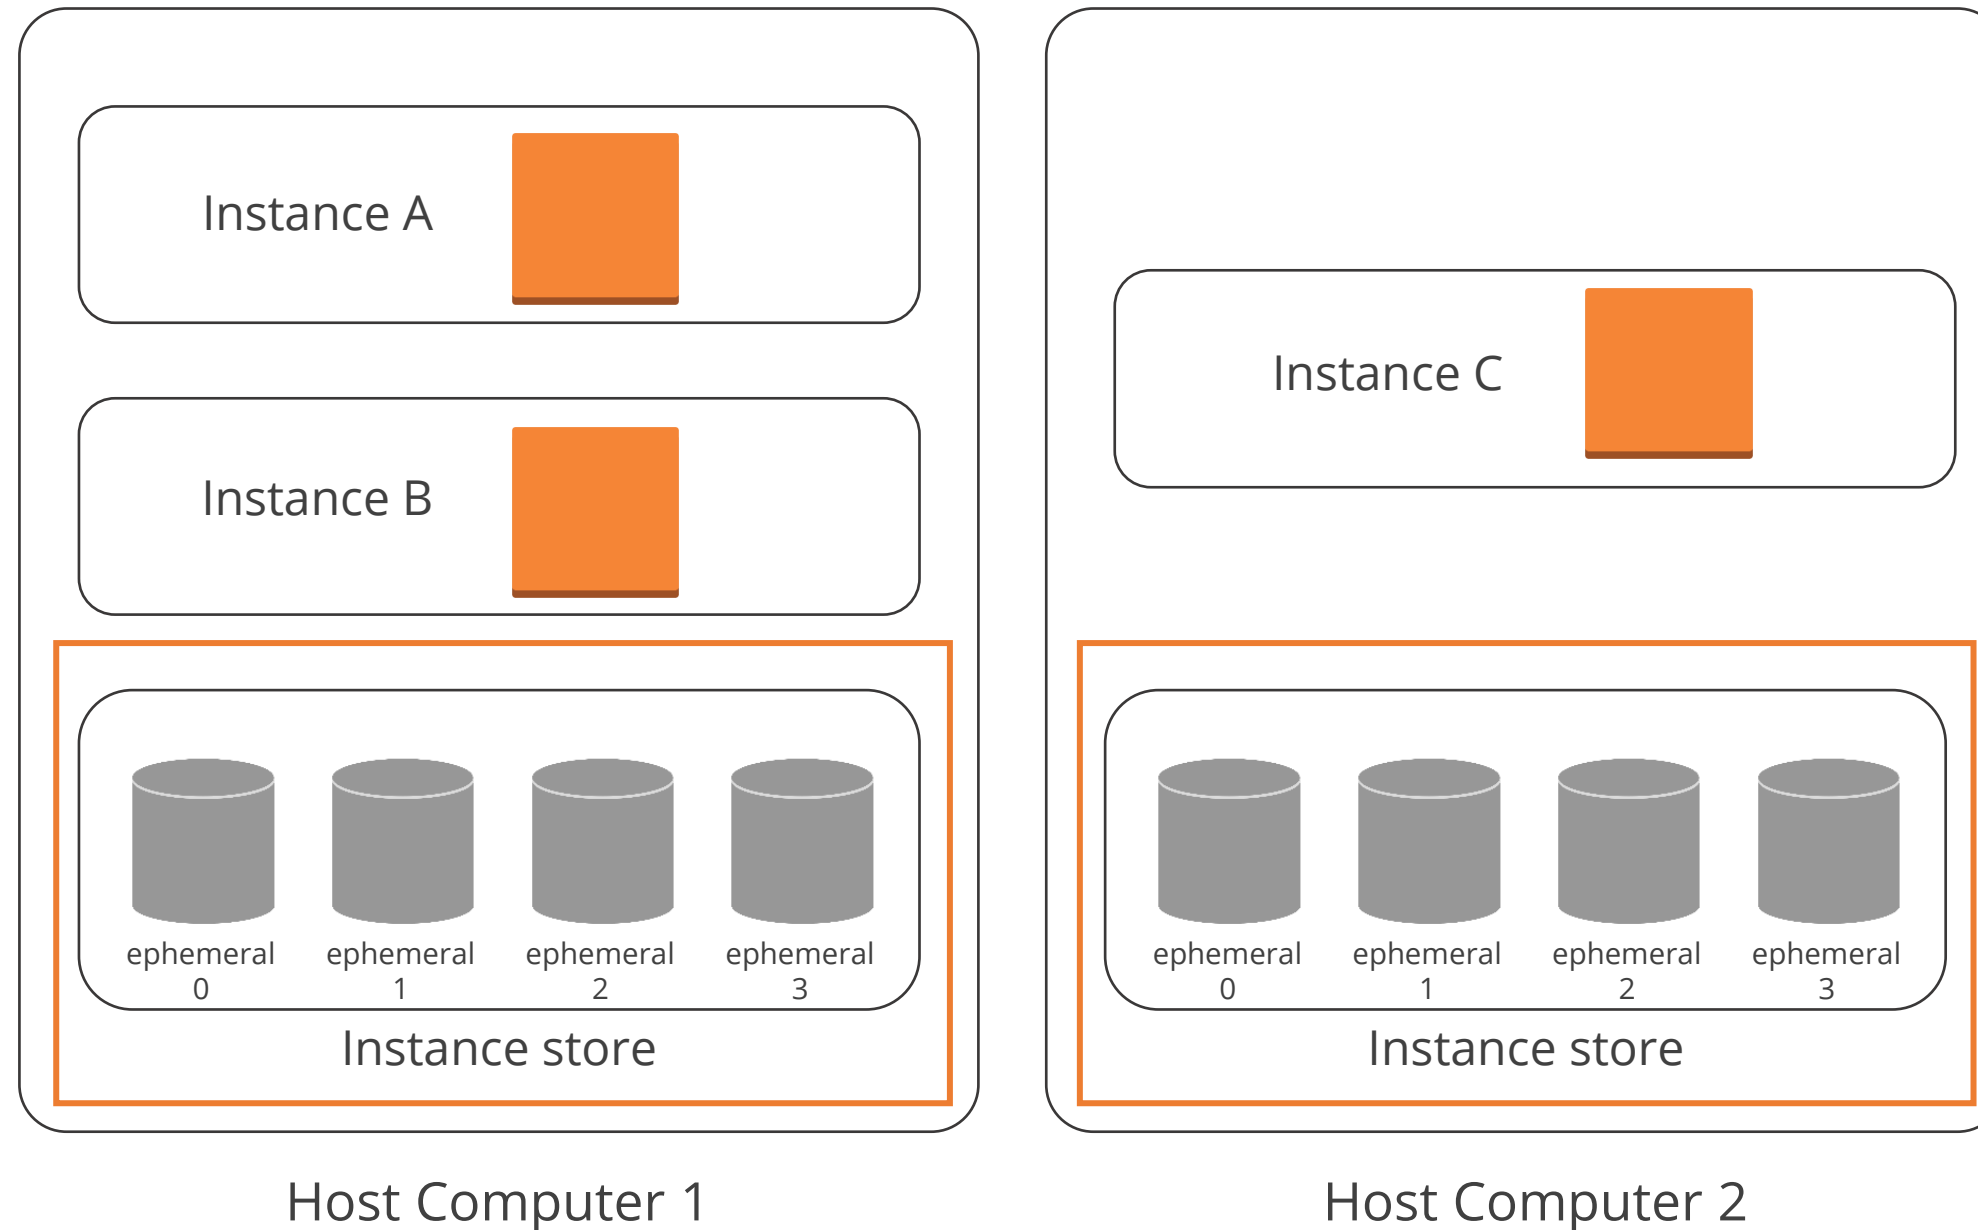


Storage Categories

	SSD-backed		HDD-backed	
Type	Provisioned IOPS SSD (io1)	General Purpose SSD (gp2)	Throughput Optimized HDD (st1)	Cold HDD (sc1)
Purpose	High performance SSD volume designed for latency-sensitive transactional workloads	General Purpose SSD volume that balances price performance for a wide variety of transactional workloads	Low-cost HDD volume designed for frequently accessed, throughput-intensive workloads	Lowest cost HDD volume designed for less frequently accessed workloads
Use Cases	I/O-intensive NoSQL and relational databases	Boot volumes, low-latency interactive apps, dev and test	Big data, data warehouses, log processing	Colder data requiring fewer scans per day
Volume Size	4 GB -16 TB	1 GB -16 TB	500 GB -16 TB	500 GB -16 TB
Max IOPS/Volume	64,000	16,000	500	250
Max Throughput/Volume	1,000 MB/s	250 MB/s	500 MB/s	250 MB/s
Max IOPS/Instance	80,000	80,000	80,000	80,000
Max Throughput/Instance	2,375 MB/s	2,375 MB/s	2,375 MB/s	2,375 MB/s
Price	\$0.125/GB-month + \$0.065/provisioned IOPS	\$0.10/GB-month	\$0.045/GB-month	\$0.025/GB-month
Dominant Performance Attribute	IOPS	IOPS	MB/s	MB/s

Instance Store

An instance store provides temporary block-level storage for your instance.



Amazon EBS vs. Instance Store

This table shows the different characteristics between the EBS-backed and Amazon Instance store-backed volumes.

Characteristics	Amazon EBS	Instance Store
Boot time	Usually less than 1 minute	Usually less than 5 minutes
Root device volume	Amazon EBS volume	Instance store volume
Upgrading	EBS-backed instances can be upgraded for instance type, Kernel, RAM disk, and user data	Instance store-backed instances cannot be upgraded
AMI creation	AMI can be easily created using a single command	AMI creation requires AMI tools and needs to be executed from within the running instance
Volume attachment	EBS volume can be attached as additional volumes when the instance is being launched and even when the instance is up and running	Instance store volume can be attached as additional volumes only when the instance is being launched and not when the instance is up and running

EBS Snapshot

Features of EBS Snapshots

- You can back up EBS volumes by taking point-in-time snapshots and it is stored in the S3 owned by AWS.
- You can perform incremental backups, that is back up only the blocks on the device that have changed after the most recent snapshot is saved.
- When you delete a snapshot, only the data exclusive to that snapshot is removed.
- Active snapshots contain all the information needed to restore your data (from the time the snapshot was taken) to a new EBS volume.
- If the EBS volume is encrypted, then the snapshot is also encrypted.



EBS Snapshot

EBS Snapshot Storage

With Amazon EBS, you can:

- Create a point-in-time snapshot of volumes that are stored for you in Amazon Simple Storage Service (Amazon S3)
- Copy a snapshot from one AWS region to another or within the same region
- Encrypt your data with Amazon S3 server-side encryption (256-bit Advanced Encryption Standard)

Assisted Practice

Create and Attach an EBS Volume to a Linux Instance

Duration: 25 min.

Problem Statement:

You are given a project to create and attach an EBS volume to a Linux Instance.

Assisted Practice: Guidelines to Create and Attach an EBS Volume to a Linux Instance

Steps to perform:

1. Open the Amazon EC2 console
2. Choose Elastic Block Store, and then Volumes from the navigation pane
3. Create an EBS Volume
4. Attach an EBS Volume to a Linux Instance
5. Convert your private key to PuTTY using PuTTYgen
6. Connect to your Linux Instance

Assisted Practice

Create a Snapshot of an Existing EBS Instance

Duration: 10 min.

Problem Statement:

You are given a project to create a snapshot of an existing EBS Instance.

Assisted Practice: Guidelines to Create a Snapshot of an Existing EBS Instance

Steps to perform:


1. Open the Amazon EC2 console
2. Choose Snapshots under Elastic Block Store in the navigation pane
3. Create Snapshot

Elastic IP Address

Elastic IP Address

An Elastic IP address is a static IPv4 address associated with the AWS account in a specific AWS region.

The users can associate an Elastic IP address with their EC2 instance at any time using:



Note

There is a default limit of five Elastic IP addresses per region.

Elastic IP Address Overview

To allocate and associate an Elastic IP address with the EC2 Windows or Linux instance, follow these steps:

01

Allocate an Elastic IP address from either Amazon's pool of public IPv4 addresses or a custom IP address pool that the users bring to their AWS account

02

Associate the Elastic IP address with a running instance

Note

The user can also disassociate an existing Elastic IP address, and then reassociate it with a different instance.

Elastic IP Address Overview

- An Elastic IP address can be associated with a single instance or network interface at a time.
- The user can move an Elastic IP address from one instance or network interface to another.
- Elastic IP addresses for IPv6 are not supported.
- An Elastic IP address is accessed through the internet gateway of a Virtual Private Cloud (VPC).
- A disassociated Elastic IP address remains allocated to the account until the user explicitly releases it.

Elastic Public IP Address Overview

The users cannot manually associate or disassociate a public IP address from their instance. AWS releases the public IP address from their instance, or assigns it a new one in the following cases:

- AWS releases the instance's public IP address when it is stopped, hibernated, or terminated. The stopped or hibernated instance receives a new public IP address when it is started.
- AWS releases the instance's public IP address when the user associates an Elastic IP address with it. When the user disassociates the Elastic IP address from the instance, it receives a new public IP address.
- If the public IP address of the instance in a VPC has been released, it will not receive a new one if there is more than one network interface attached to the instance.
- If the instance's public IP address is released while it has a secondary private IP address that is associated with an Elastic IP address, the instance does not receive a new public IP address.

EC2 Instance Metadata

EC2 Instance Metadata

- Instance metadata is data about your instance that you can use to configure or manage the running instance.
- Instance metadata is divided into categories such as hostname, events, and security groups.
- You can also use instance metadata to access user data that you specified while launching your instance.

Note

- Metadata is the information about the EC2 instance.
- User data is the launch script of the EC2 instance.

EC2 Instance Metadata

You can access instance metadata from a running instance using one of the following methods:

- Instance Metadata Service Version 1 (IMDSv1) – a request/response method
- Instance Metadata Service Version 2 (IMDSv2) – a session-oriented method

To view all the categories of instance metadata from within a running instance, use the following URL:

IMDSv1	IMDSv2
<code>curl http://169.254.169.254/latest/meta-data/</code>	<code>TOKEN=`curl -X PUT "http://169.254.169.254/latest/api/token" -H "X-aws-ec2-metadata-token-ttl-seconds: 21600" `\ && curl -H "X-aws-ec2-metadata-token: \$TOKEN" -v http://169.254.169.254/latest/meta-data/</code>

Instance Metadata Categories

The different categories of instance metadata are:

- ami-id
- ami-launch-index
- ami-manifest-path
- block-device mapping/
- hostname
- iam/
- instance-action
- instance-id
- instance-type
- local-hostname
- local-ipv4
- mac
- metrics/
- network/
- placement/
- profile
- public-hostname
- public-ipv4
- public-keys/
- reservation-id
- security-groups
- services/

Elastic File System

Elastic File System (EFS)

Amazon Elastic File System (Amazon EFS) provides a simple, scalable, fully manageable elastic NFS file system for use with AWS Cloud services and on-premises resources.



Amazon EFS

Elastic File System (EFS) Overview

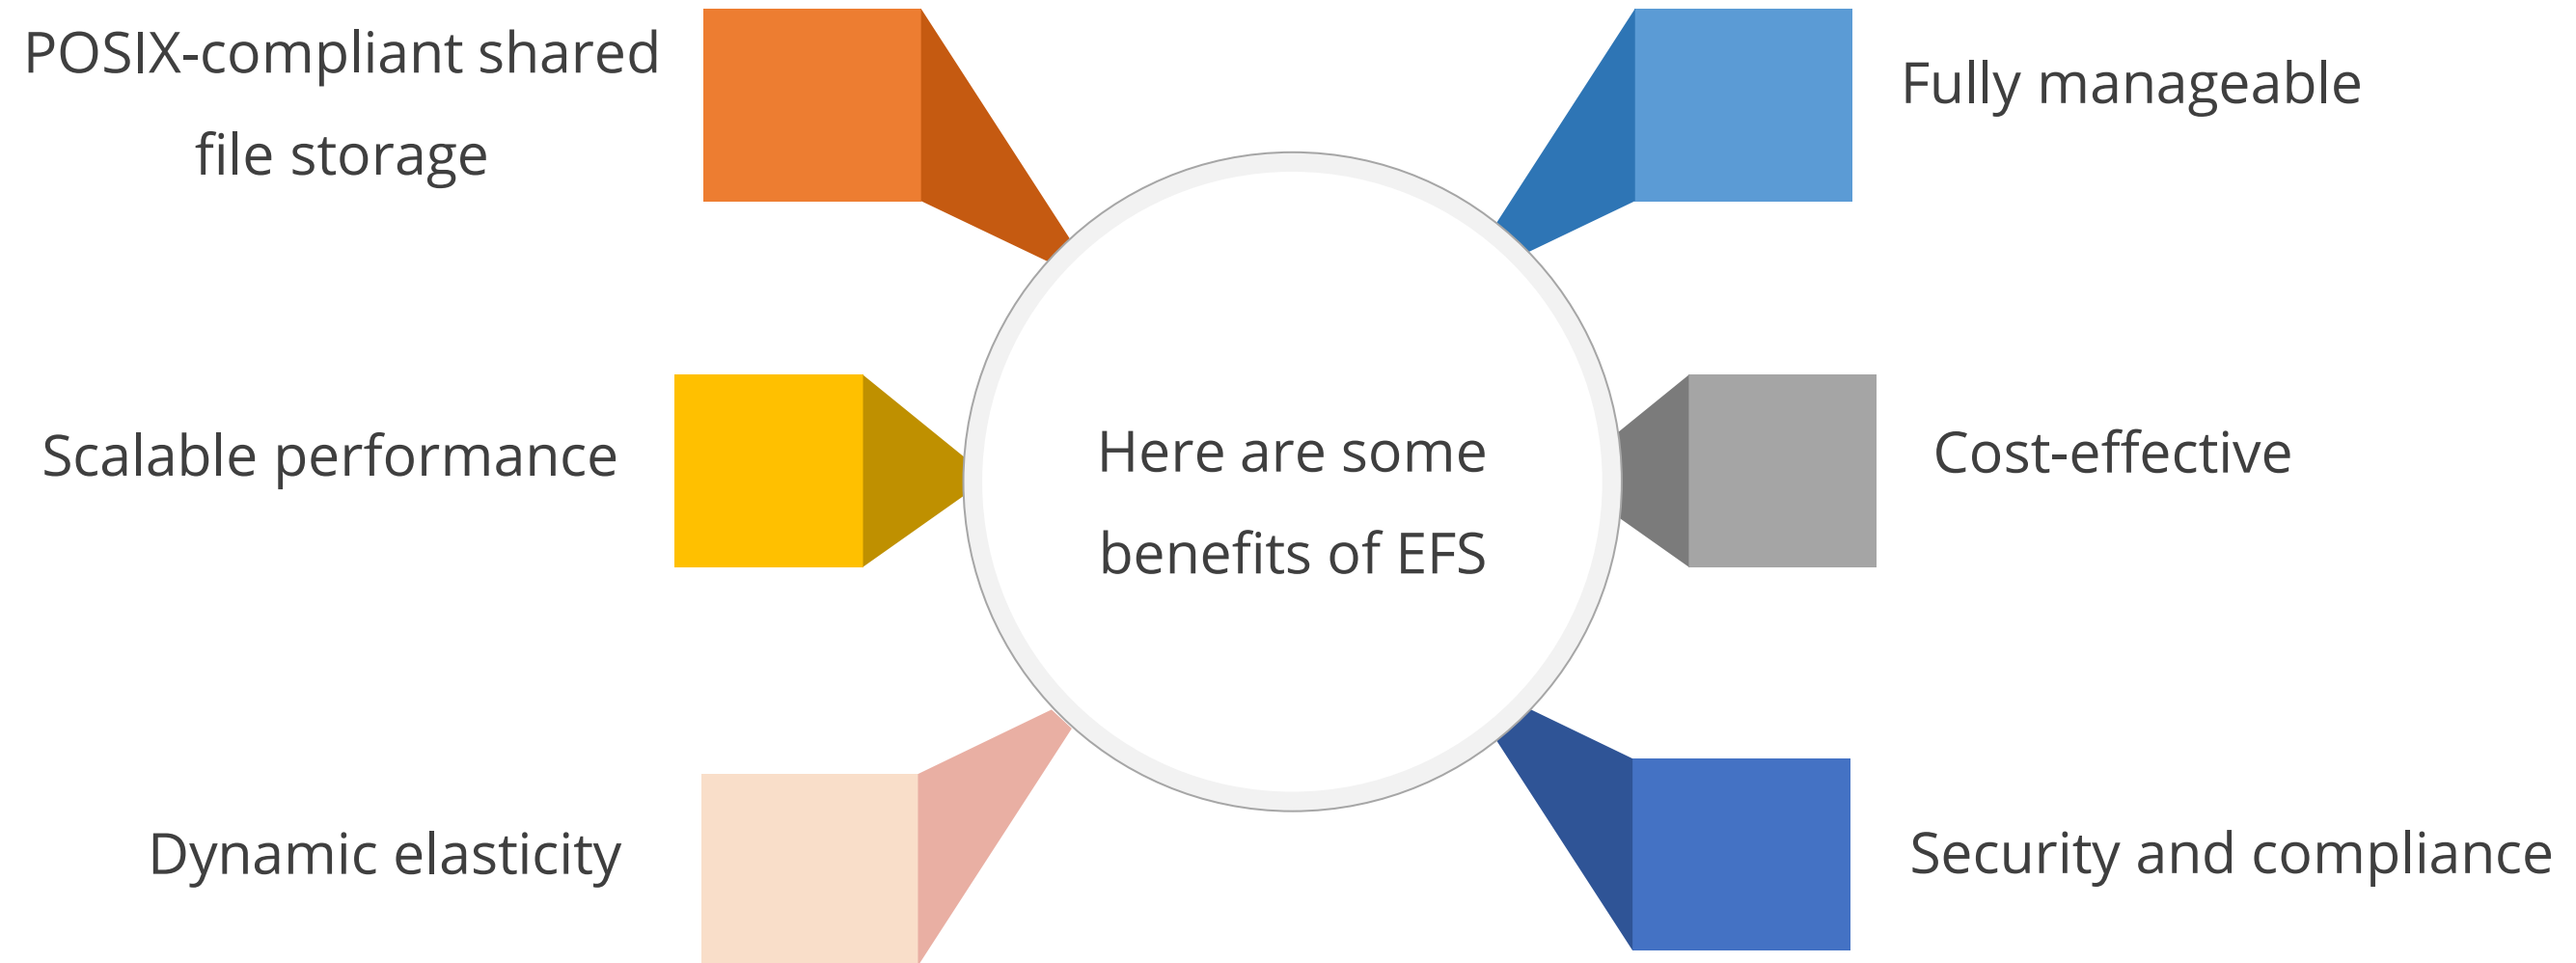
EFS supports Network File System Version 4.0 and 4.1 protocol (NFSv4).

With Amazon EFS storage, you pay only for what you use.



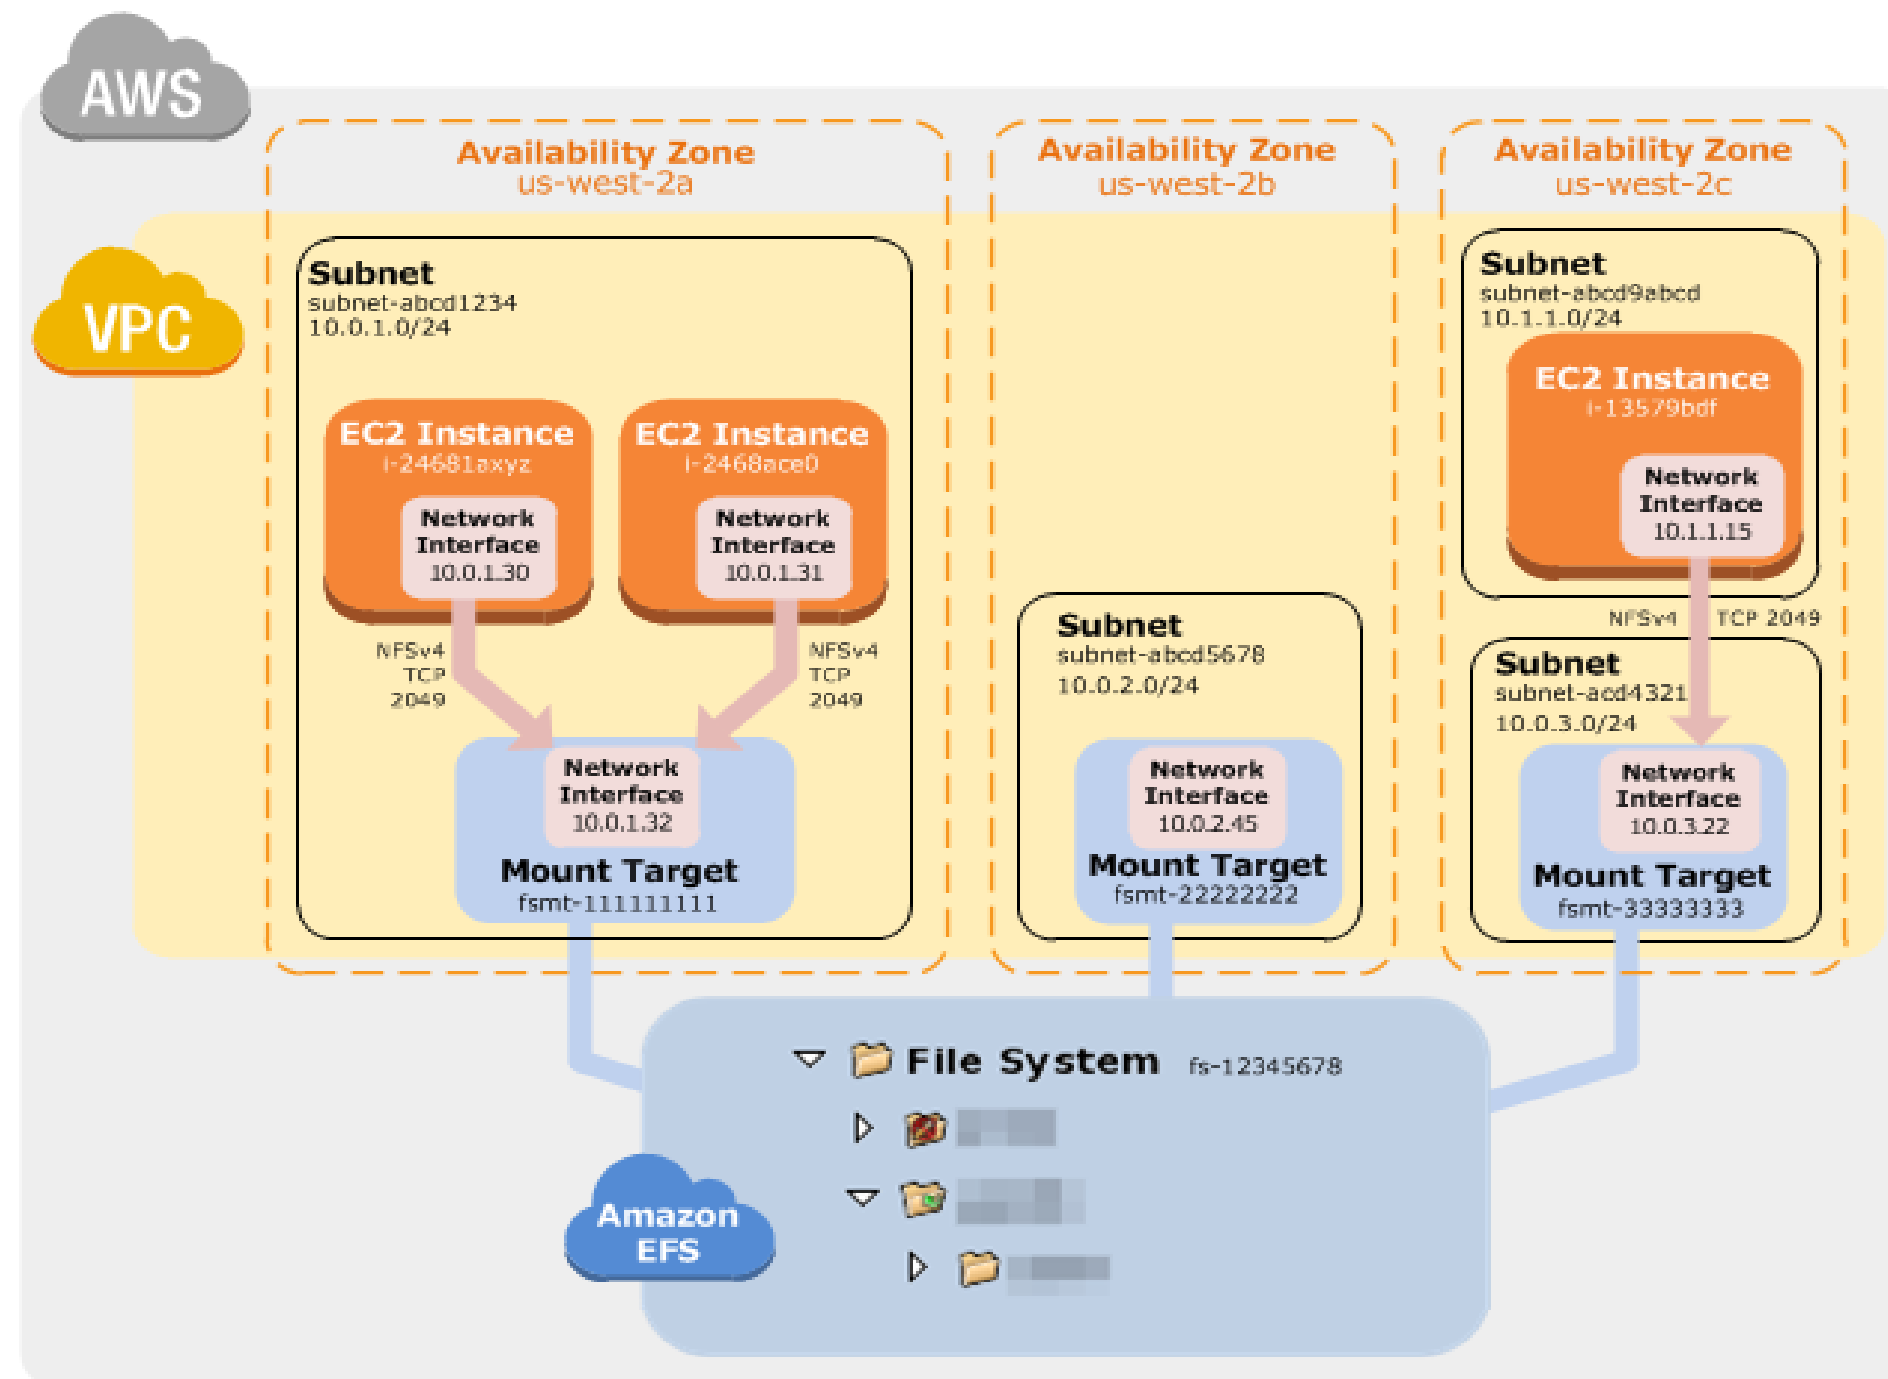
With EFS, the data is stored and accessed from all the Availability Zones in the AWS region.

Benefits of Elastic File System (EFS)



How Amazon EFS Works with Amazon EC2

The following illustration shows an example of VPC accessing an Amazon EFS file system:



Difference between EFS and EBS

EFS vs. EBS

Feature	EFS	EBS
Storage Size	No limitations	Maximum 16 TiB
Storage Type	Object storage	Block storage
Performance	Scalable	Hardly scalable
File Size Limitation	Maximum file size 47.9 TiB	No limitation
Data Throughput	Default throughput of 3 GB	SSD- and HDD-backed storage types
Data Access	Can be accessed concurrently	Limited to single EC2 instance
Availability Zone (AZ) Failure	Can survive one AZ failure	Cannot withstand AZ failure without snapshots

EC2 Instance Pricing

EC2 Instance Pricing

The following are the EC2 Instance pricing:

01	On-Demand
02	Spot
03	Reserved
04	Saving Plans
05	Dedicated

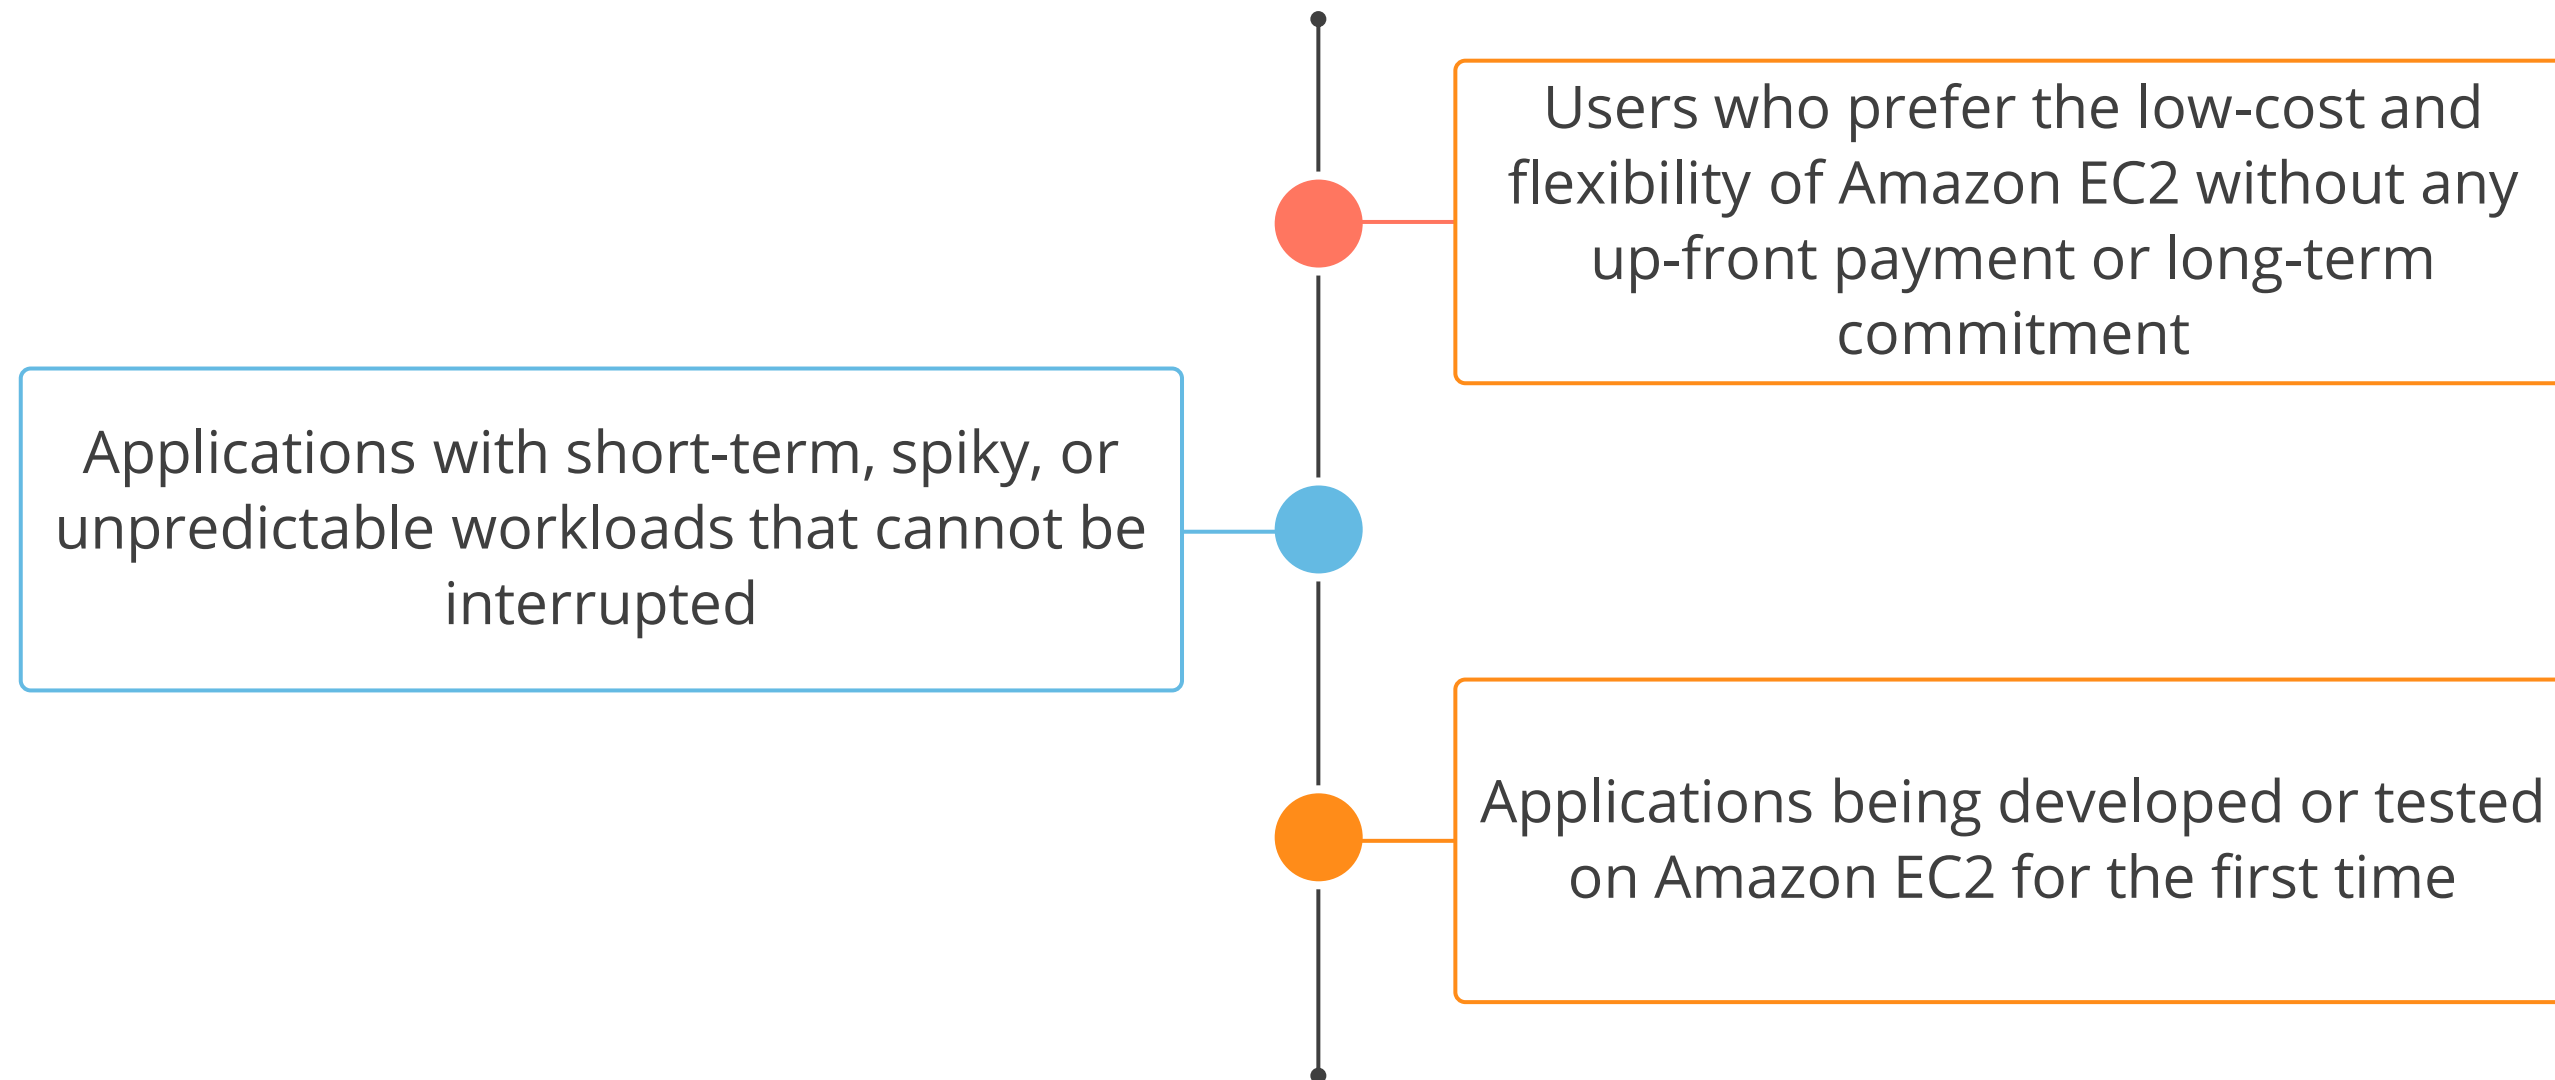
Purchasing Options: On-Demand

EC2 On-Demand instance pricing enables you to pay only for what you use with no long-term commitments.



Purchasing Options: On-Demand

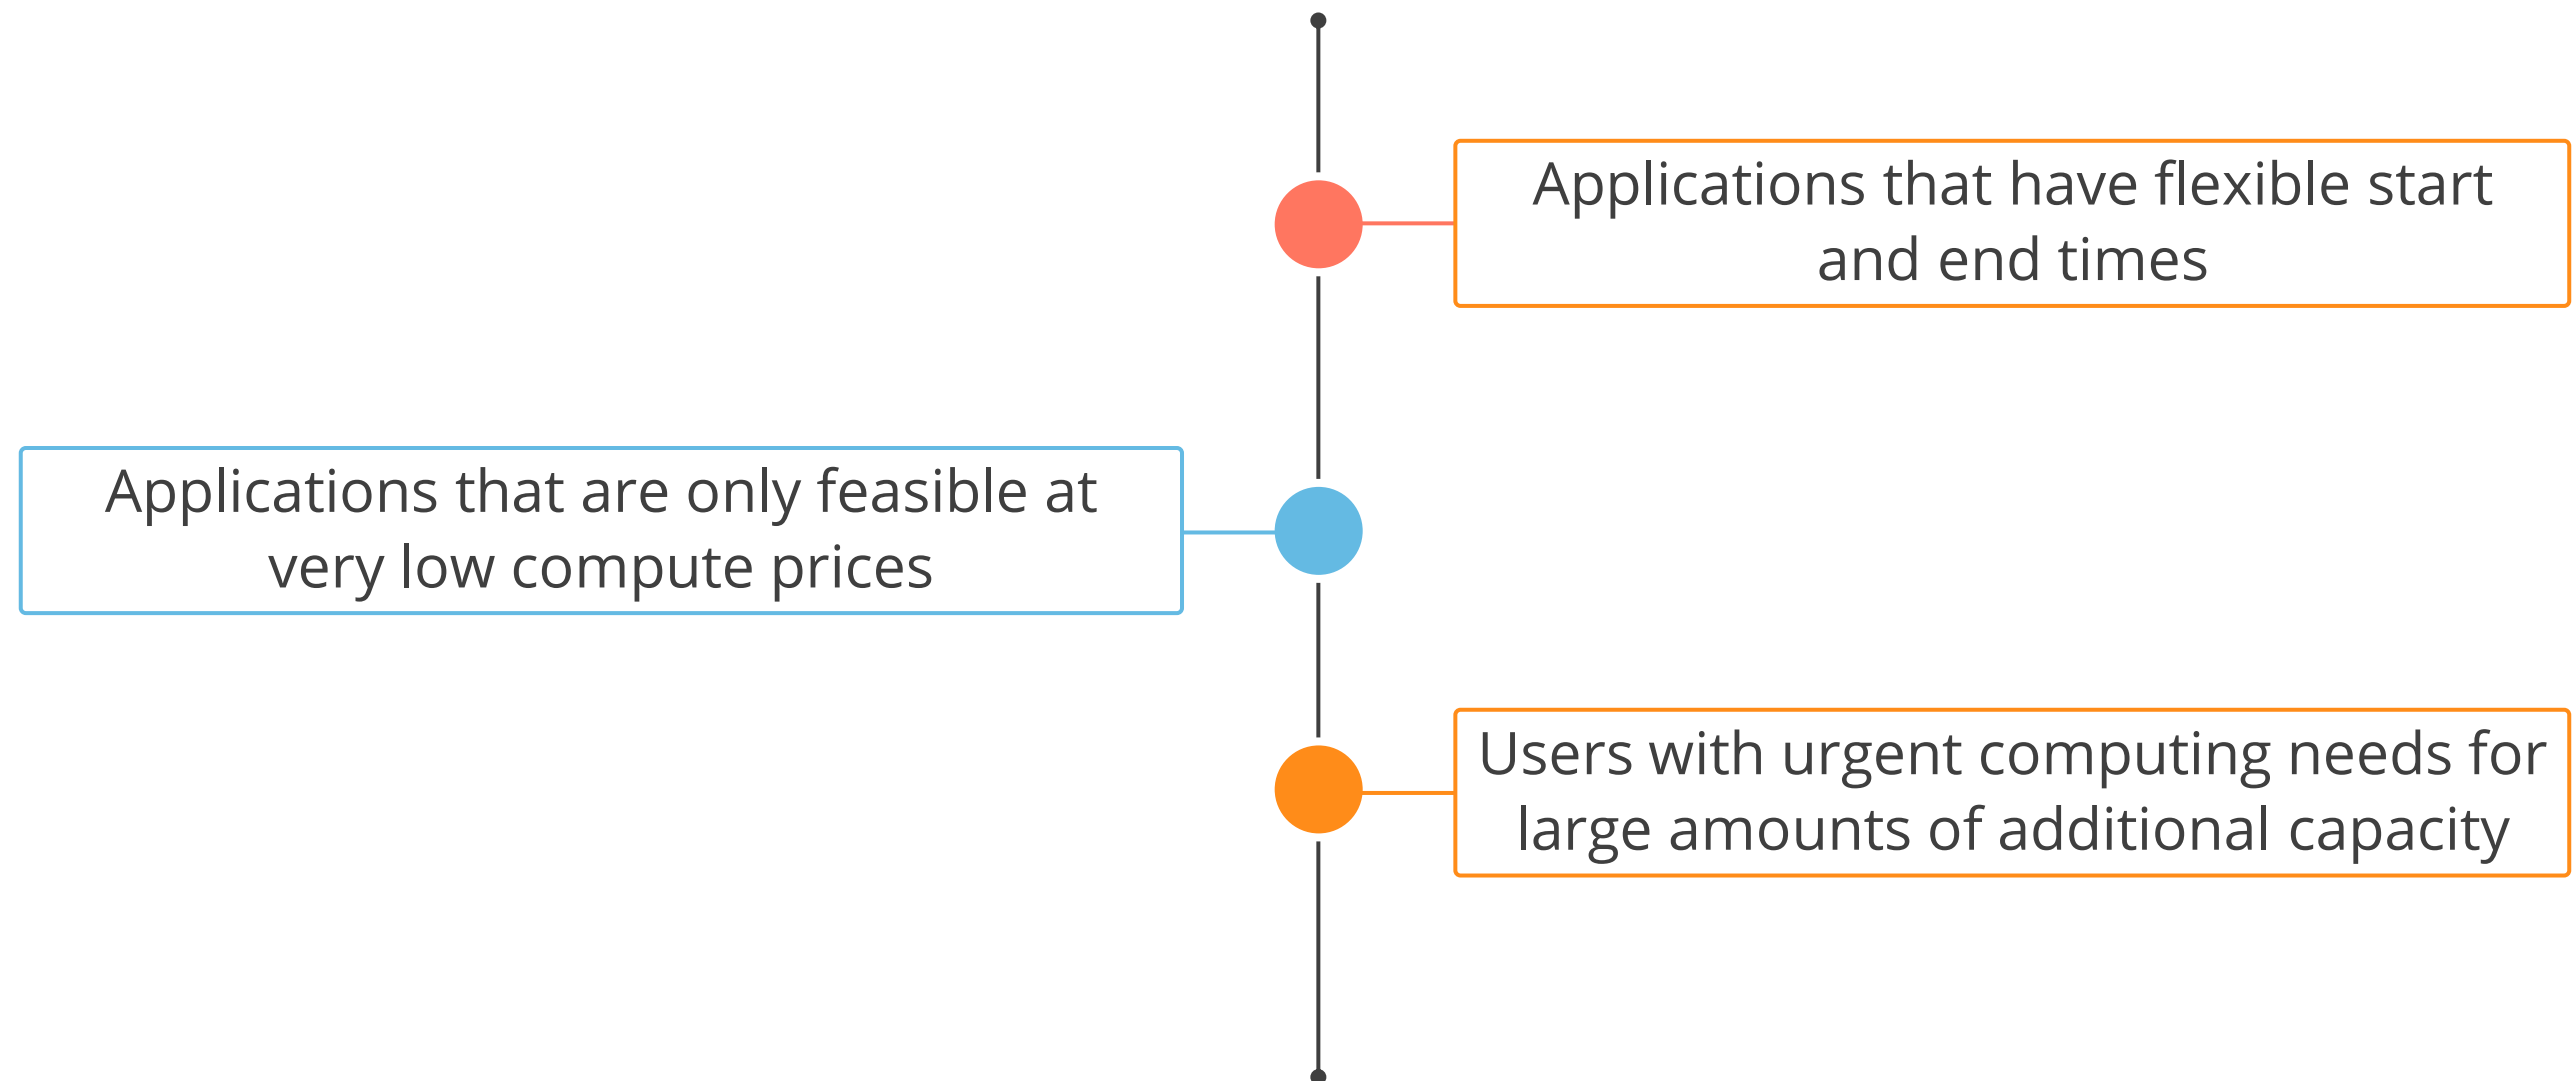
On-Demand Instances are recommended for:



Purchasing Options: Spot

Amazon EC2 Spot Instances allow you to request spare computing capacity for up to 90% off the On-Demand price.

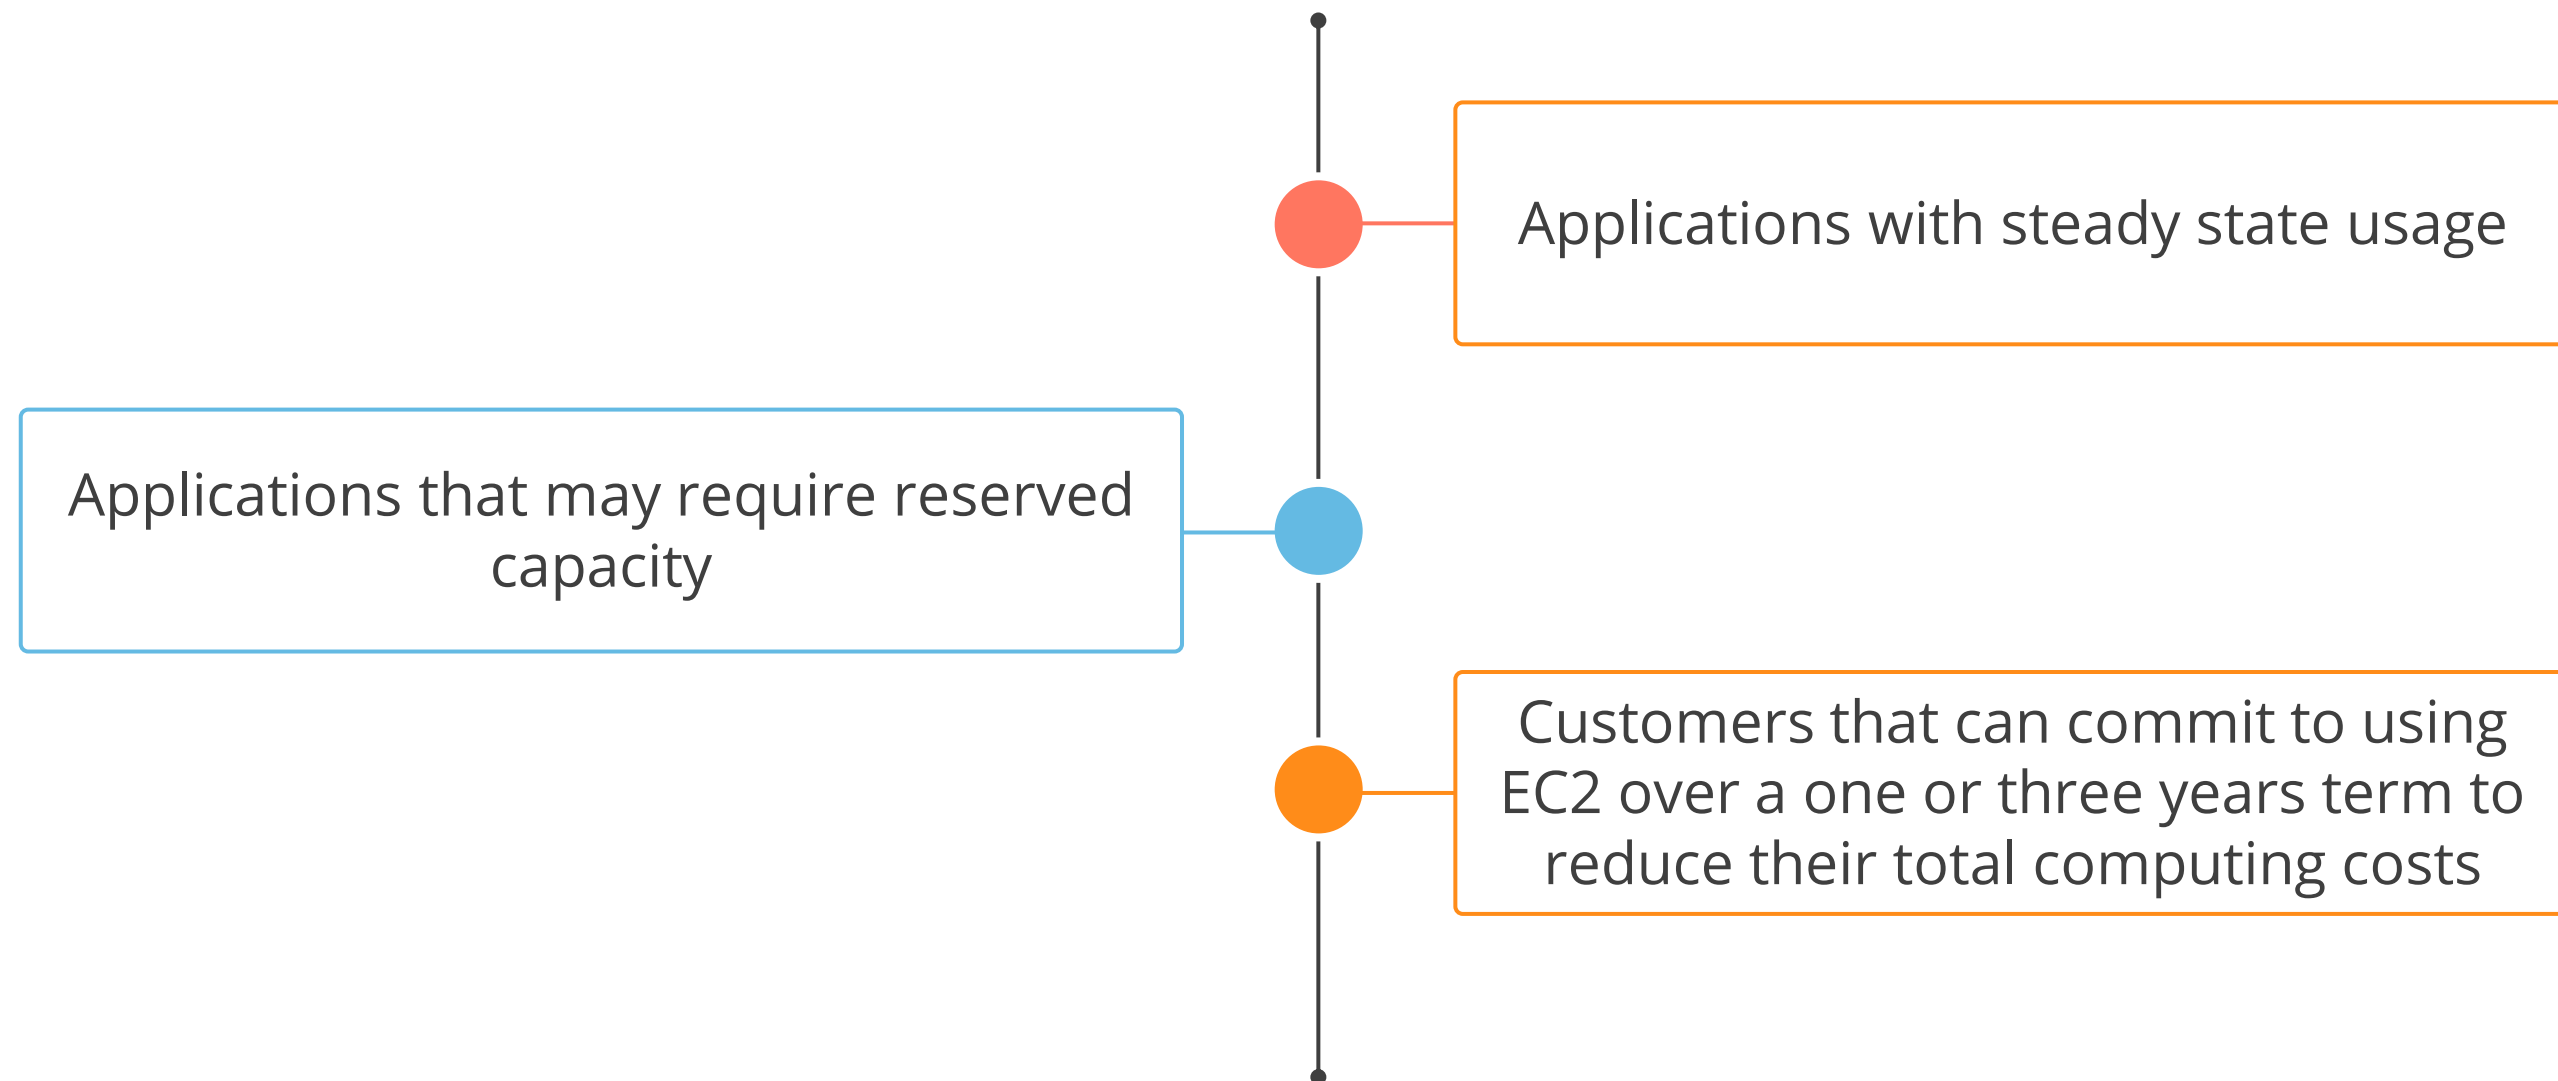
Spot Instances are recommended for:



Purchasing Options: Reserved

Reserved Instances provide you with a significant discount (up to 75%) compared to On-Demand instance pricing.

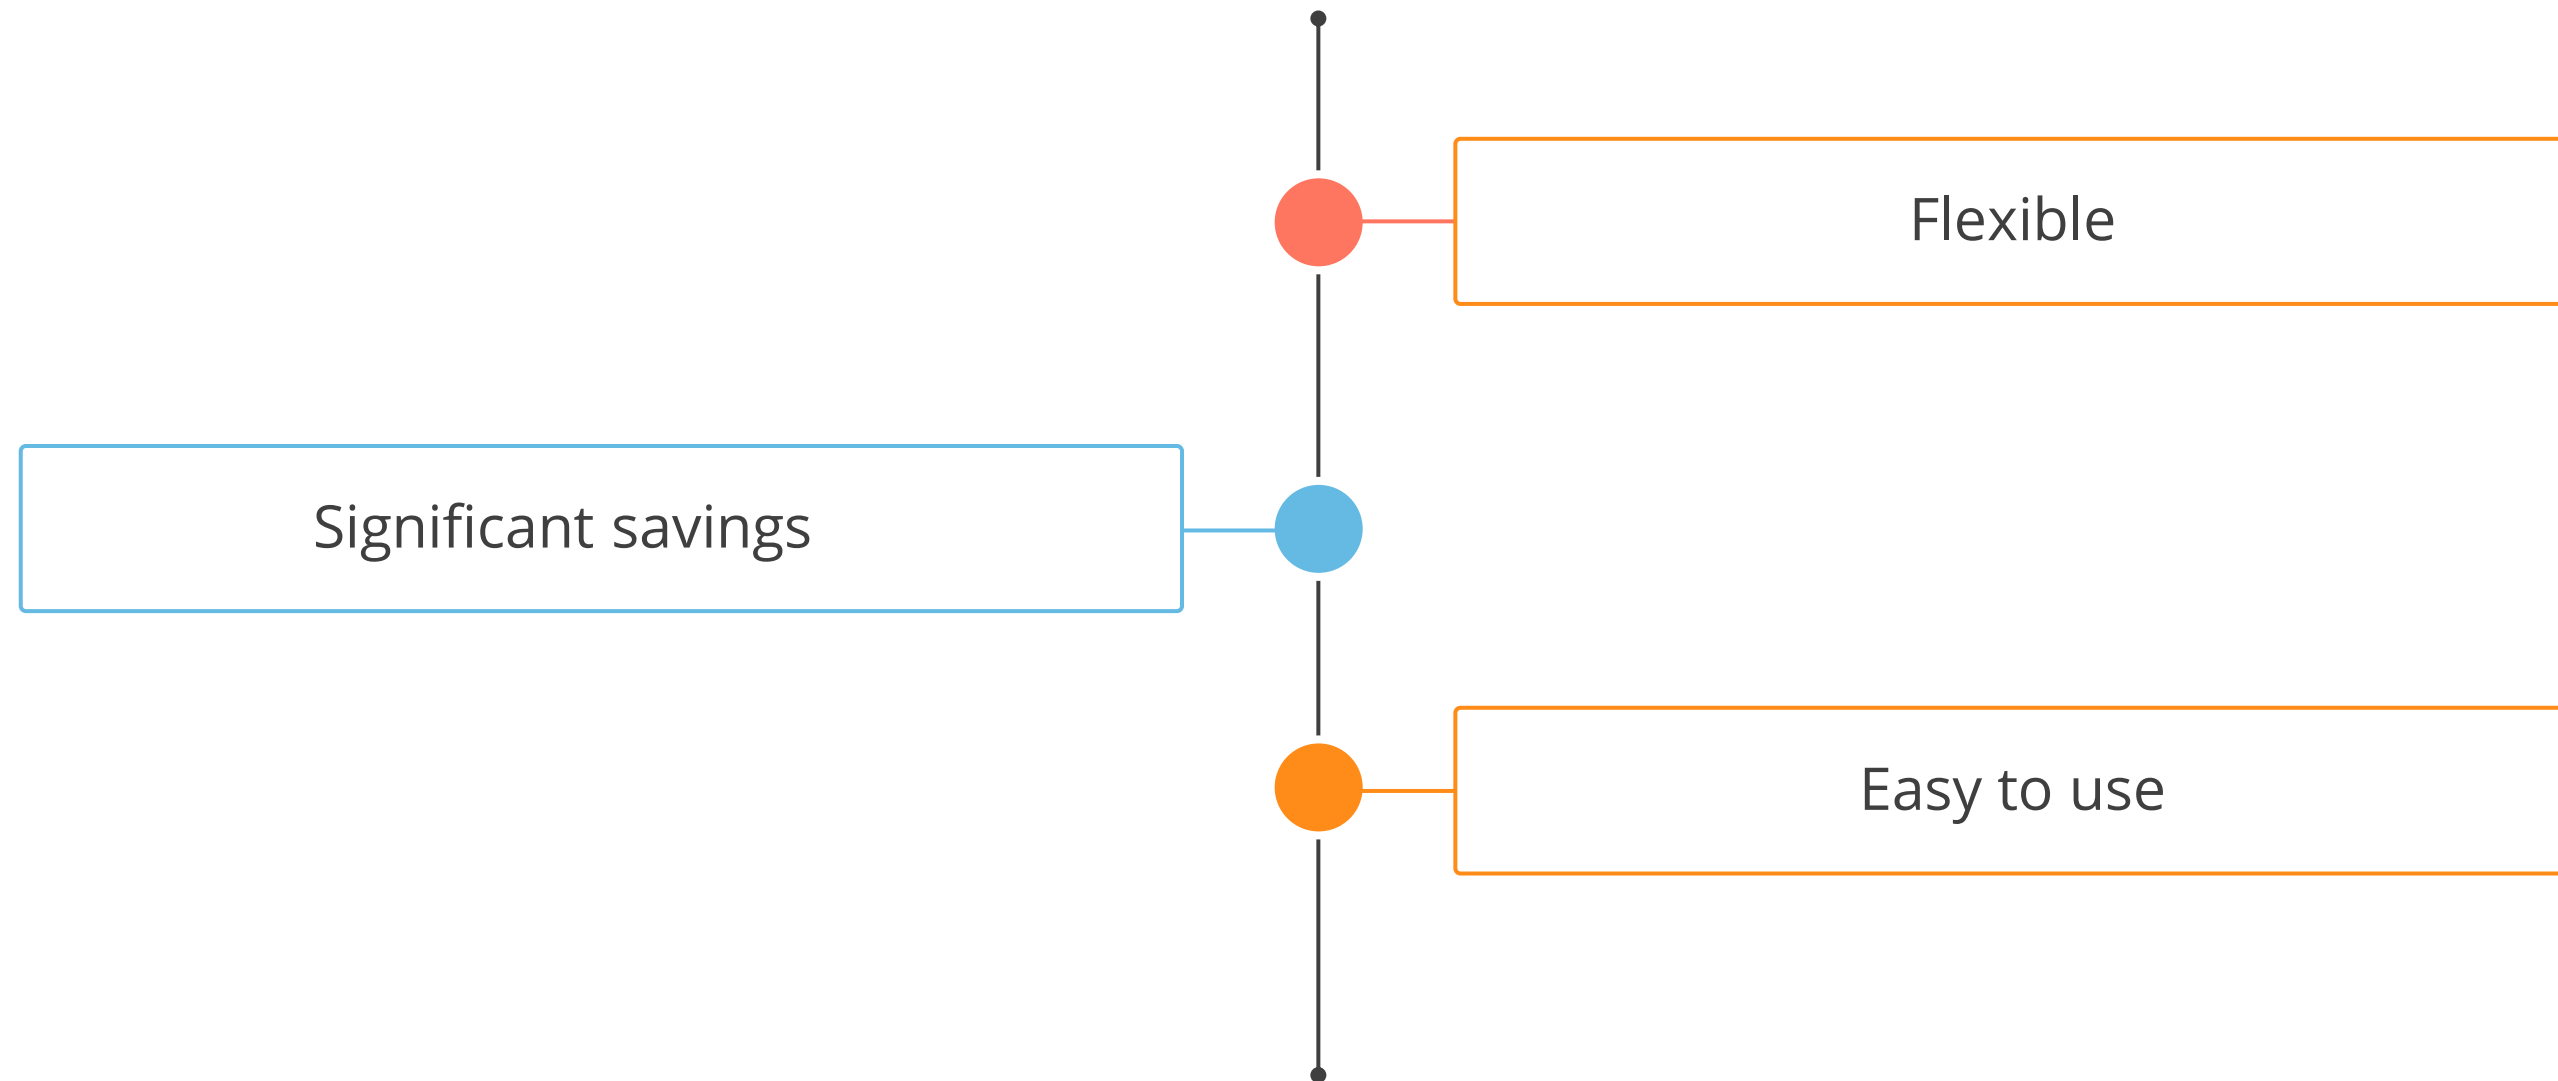
Reserved Instances are recommended for:



Purchasing Options: Savings Plans

Savings Plans are a flexible pricing model that offers low prices on EC2 usage, in exchange for a commitment to a consistent amount of usage (measured in \$/hour), for a one or three year term.

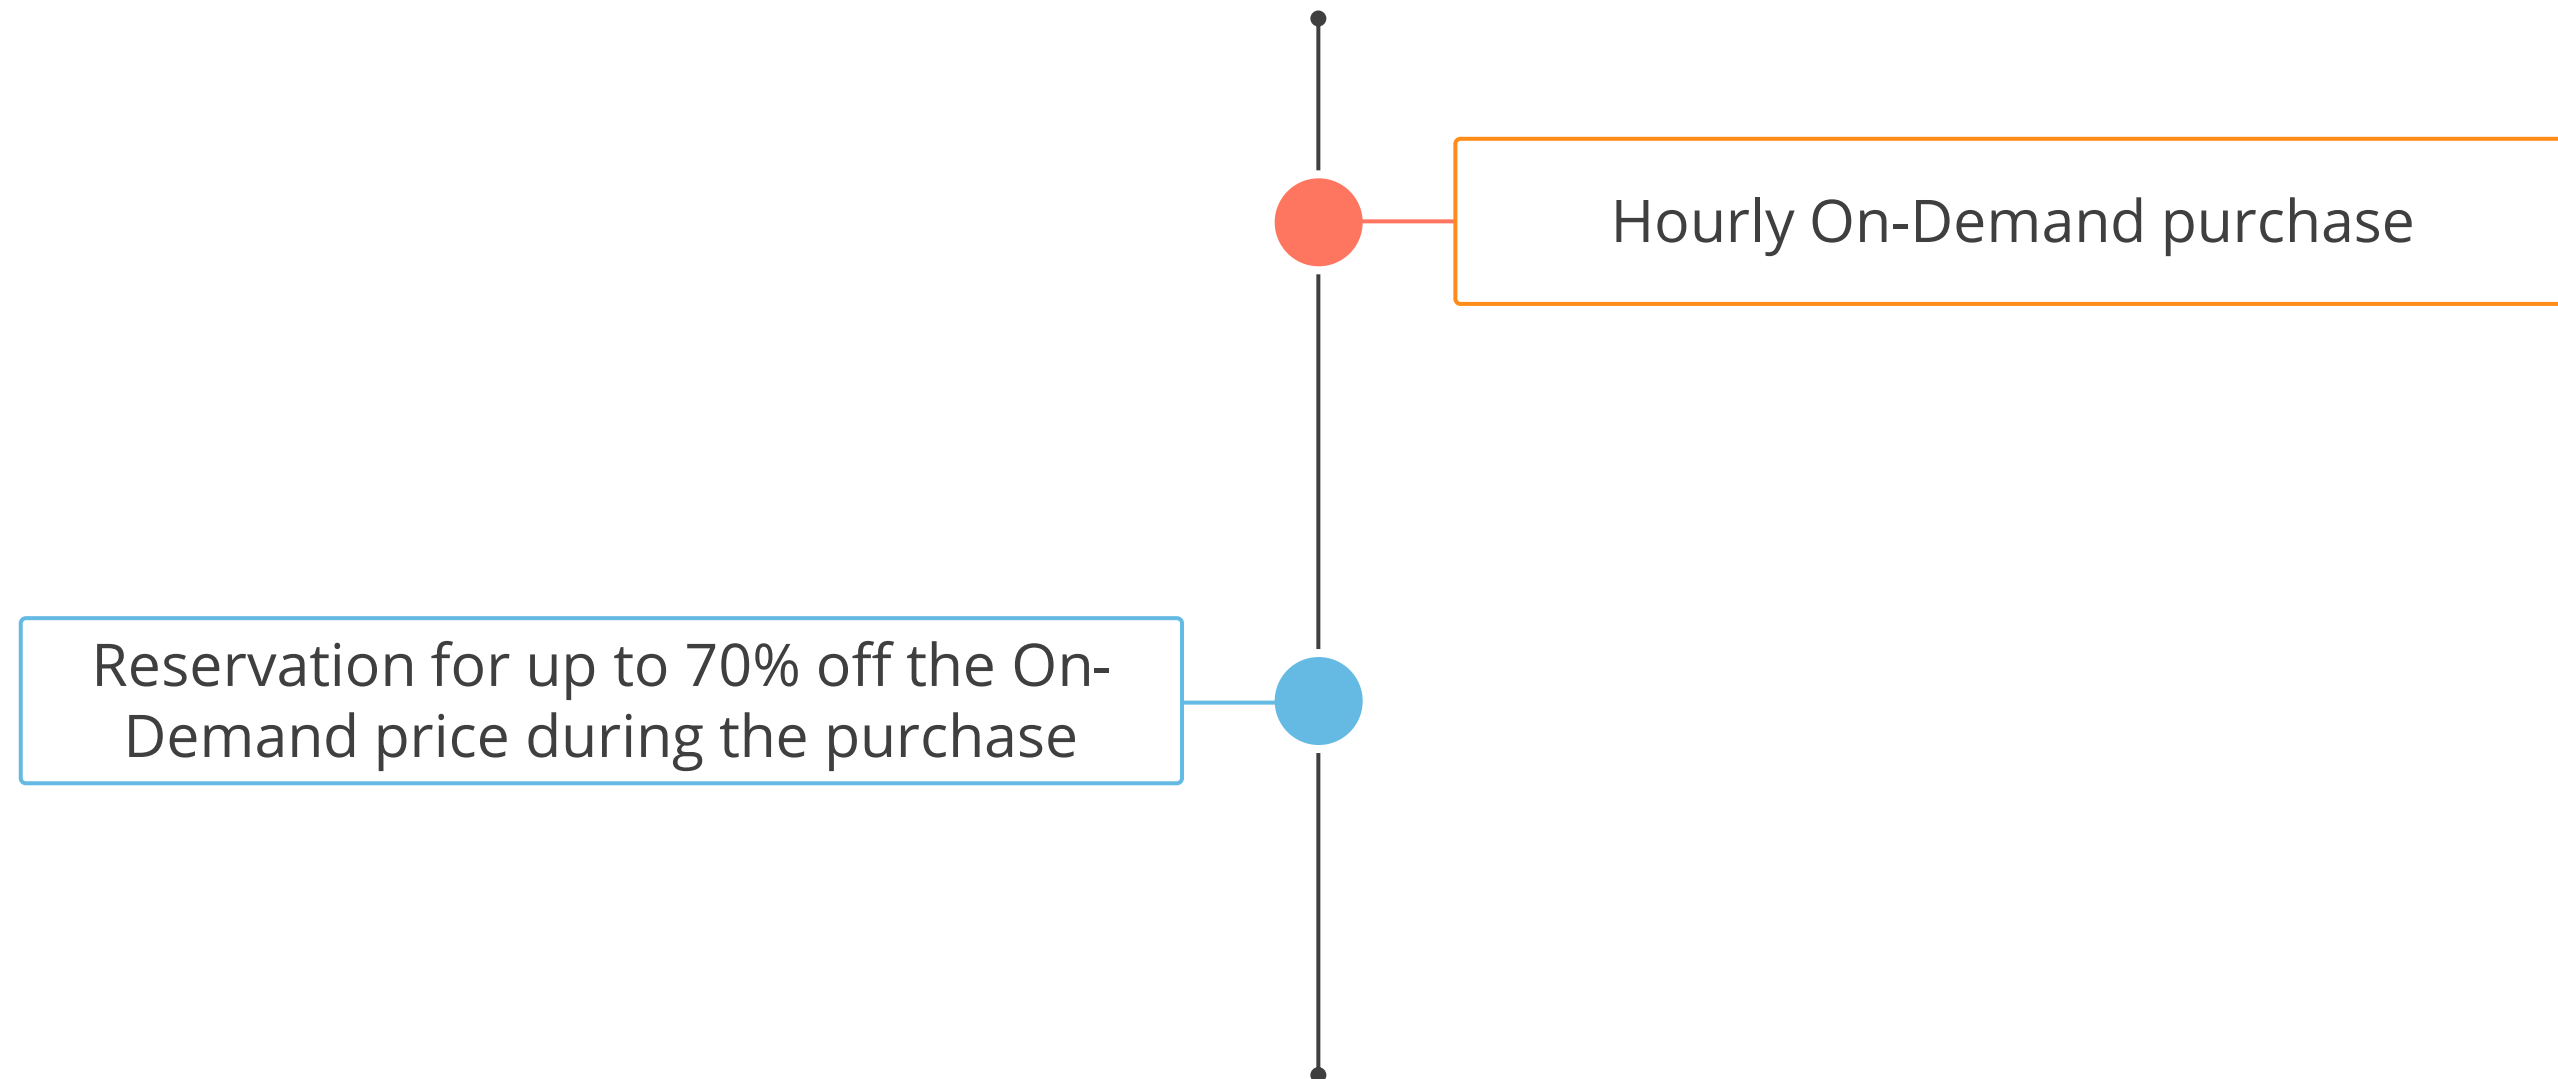
Benefits of Savings Plans are:



Purchasing Options: Dedicated

An Amazon EC2 Dedicated Host is a physical server with EC2 instance capacity fully dedicated to your use.

Dedicated Instances are recommended for:



AWS EC2 Best Practices

AWS EC2 Best Practices

Storage

- Understand the implications of the root device type for data persistence, backup, and recovery

Resource Management

- Ensure that the volume with your data persists after instance termination

Backup and Recovery

- Use the instance store to store temporary data
- Encrypt EBS volumes and snapshots

AWS EC2 Best Practices

Storage

Resource Management

Backup and Recovery

- Use instance metadata and custom resource tags to track and identify your AWS resources
- View your current limits for Amazon EC2

AWS EC2 Best Practices

Storage

Resource Management

Backup and Recovery

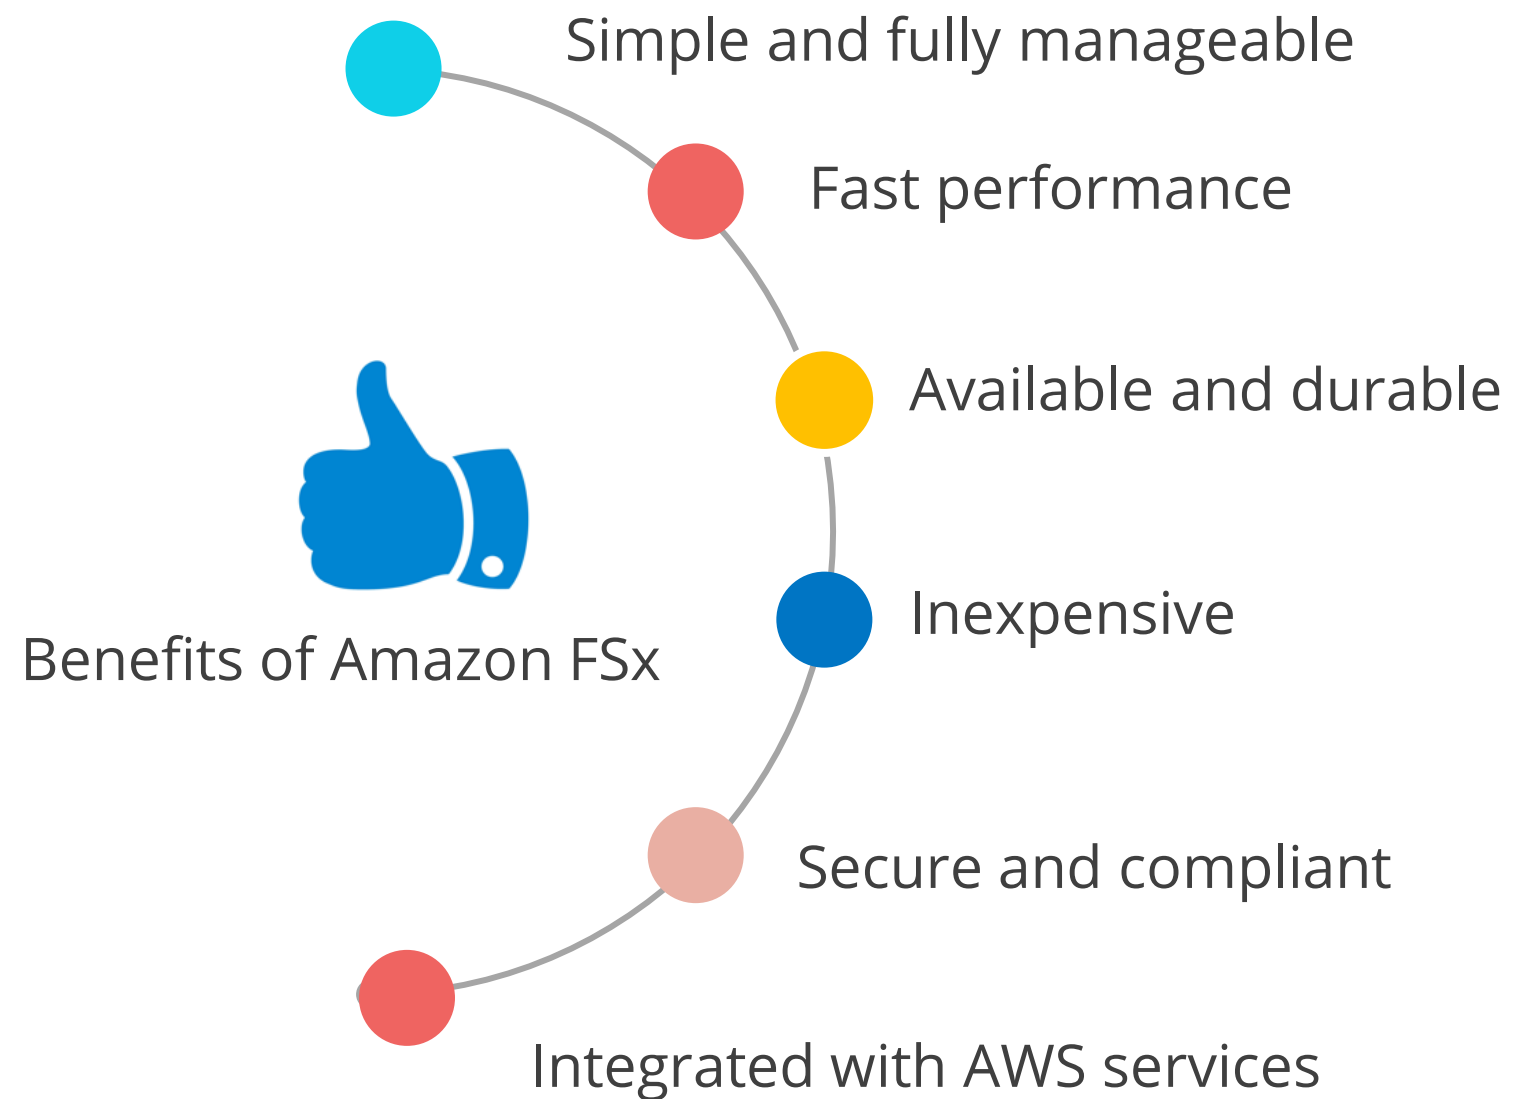
- Regularly back up your instance using Amazon EBS snapshots or a backup tool
- Deploy critical components of your application across multiple Availability Zones and replicate your data appropriately
- Design your applications to handle dynamic IP addressing while restarting your instance
- Monitor and respond to events
- Ensure you are prepared to handle failover
- Regularly test the process of recovering your instances and Amazon EBS volumes in case of failure

Amazon FSx

Amazon FSx Overview

- Amazon FSx makes it easy and cost-effective to launch and run popular file systems.
- With Amazon FSx, you can leverage the rich feature sets and fast performance of widely-used open-source and commercially-licensed file systems.
- It provides cost-efficient capacity and high levels of reliability.
- It integrates with other AWS services so that you can manage and use the file systems in cloud-native ways.

Amazon FSx Overview



Types of Amazon FSx

Amazon FSx provides you with two file systems to choose from:



Amazon FSx for Windows File Server
for business applications




Amazon FSx for Lustre for high-
performance workloads

Amazon FSx for Windows

- Amazon FSx for Windows File Server provides fully managed, highly reliable, and scalable file storage that is accessible over the industry-standard Server Message Block (SMB) protocol.
- It is built on Windows Server, delivering a wide range of administrative features such as user quotas, end-user file restore, and Microsoft Active Directory (AD) integration.
- It offers single-AZ and multi-AZ deployment options, fully managed backups, and encryption of data at rest and in transit.



Amazon FSx for Windows

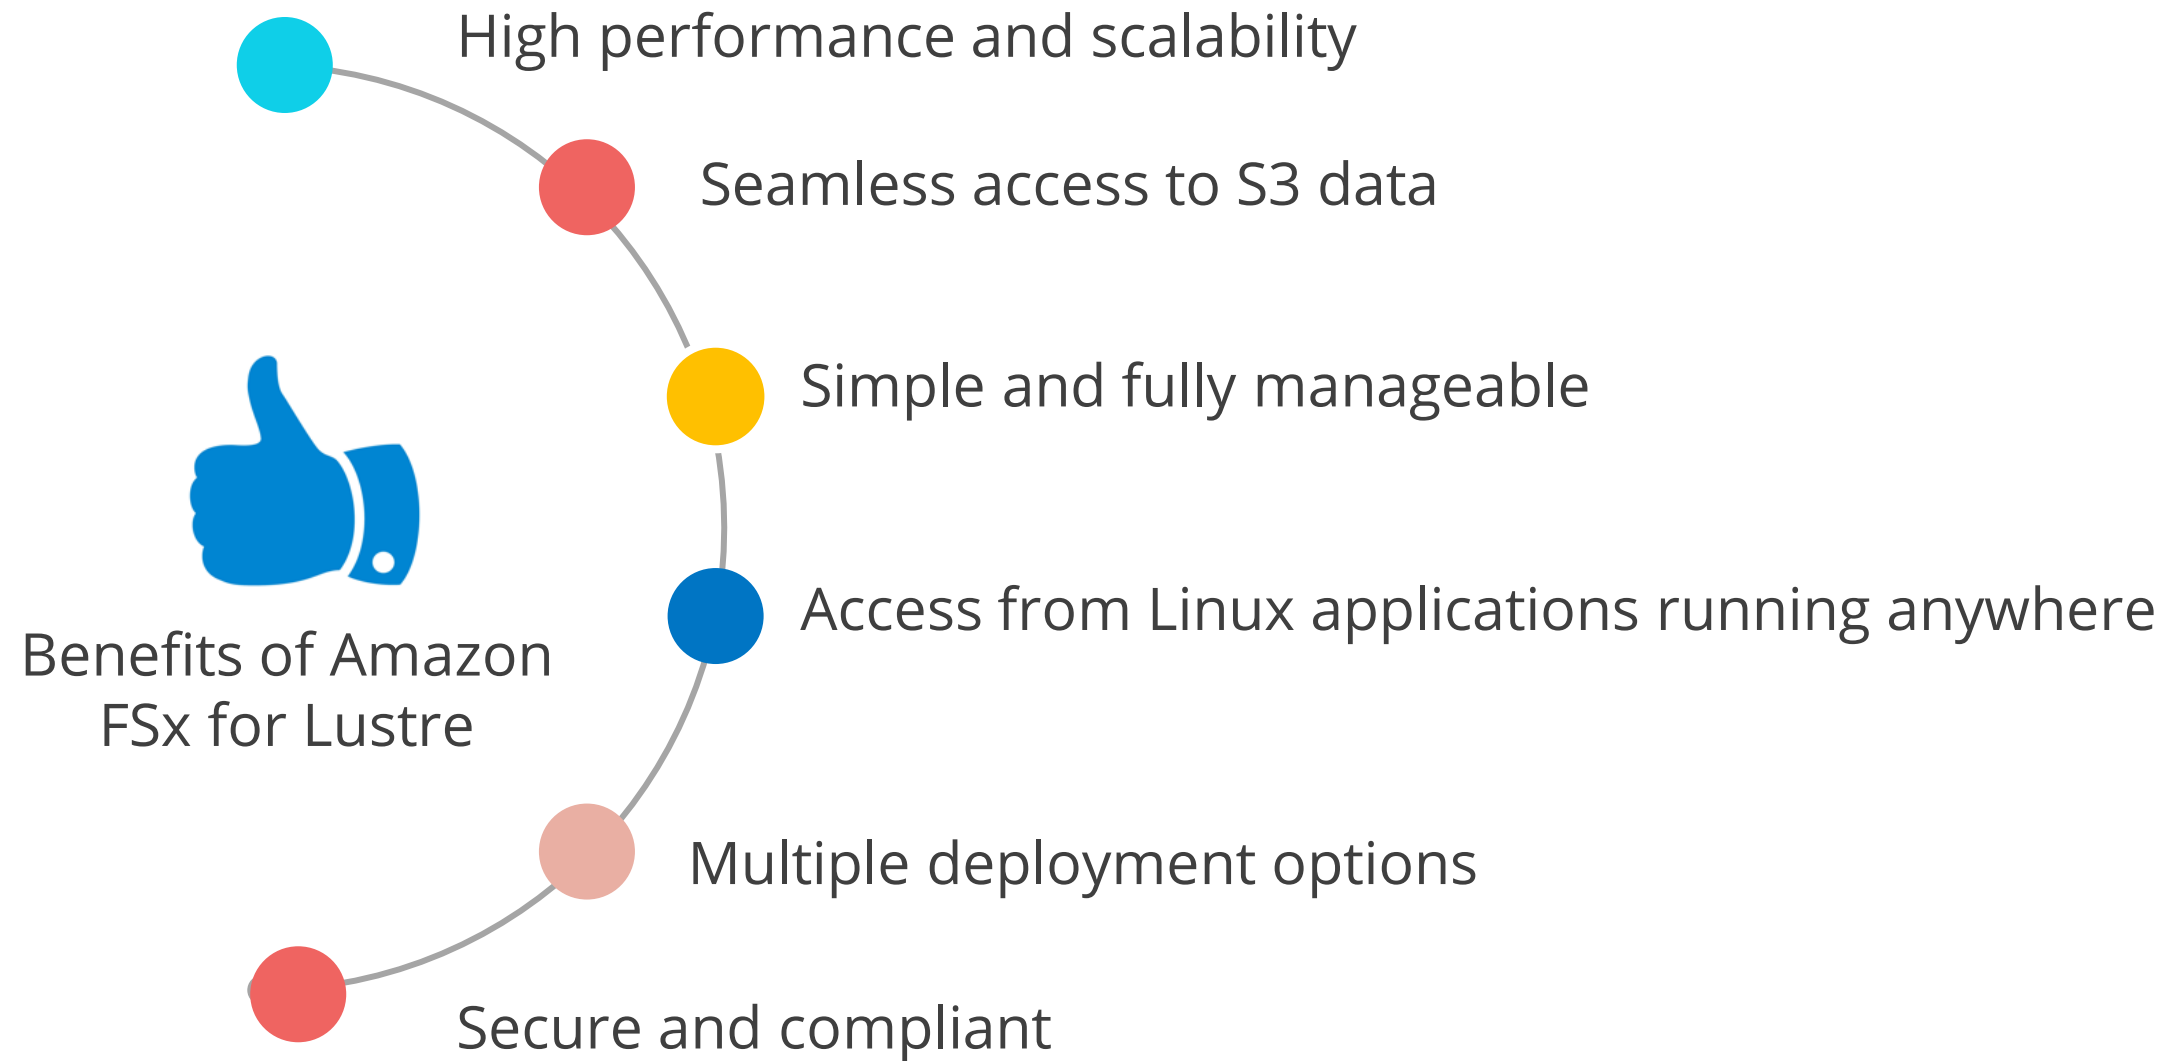


Amazon FSx for Lustre

- Amazon FSx for Lustre makes it easy and cost-effective to launch and run the world's most popular high-performance file system.
- The open-source Lustre file system is designed for applications that require fast storage that can keep up with your computing performance.
- Amazon FSx enables you to use Lustre file systems for any workload where storage speed matters.
- Amazon FSx for Lustre integrates with Amazon S3, making it easy to use Lustre file system to access data sets.



Amazon FSx for Lustre



High Performance Computing (HPC)

High Performance Computing (HPC)

- AWS provides the most elastic and scalable cloud infrastructure to run your HPC applications.
- HPC on AWS removes the long wait times and lost productivity often associated with on-premises HPC clusters.
- The most compute-intensive workloads span the traditional HPC applications, like genomics, computational chemistry, financial risk modeling, computer-aided engineering, weather prediction, and seismic imaging, as well as emerging applications, like machine learning, deep learning, and autonomous driving.
- With access to a broad portfolio of cloud-based services like Data Analytics, Artificial Intelligence (AI), and Machine Learning (ML), you can redefine traditional HPC workflows to innovate faster.

HPC Benefits

Faster results

Flexible configurations

Cost-effective operations

HPC on AWS eliminates the wait times and long job queues often associated with limited on-premises HPC resources, helping you to get results faster.

HPC Benefits

Faster results

Flexible configurations

Cost-effective operations

Moving HPC workloads to the cloud can help increase productivity by matching the infrastructure configuration to the application.

HPC Benefits

Faster results

Flexible configurations

Cost-effective operations

There is no initial capital expenditure or long procurement periods for HPC on AWS. You pay only for the storage capacity you use, and AWS offers you with flexible pricing models which help you in significant cost savings when you process time-flexible, stateless workloads.

Key Takeaways

- Amazon Elastic Compute Cloud (Amazon EC2) is a web service that provides scalable computing capacity in the Amazon Web Services (AWS) cloud.
- An AMI is a virtual instance that includes a template, launch permissions, and a block device mapping.
- The storage categories of EBS are SSD-backed and HDD-backed.
- You can back up your EBS volumes by taking point-in-time snapshots and storing them on Amazon S3.
- EC2 On-Demand Instance pricing enables you to pay only for what you use with no long-term commitments.



Setting up an Instance with EBS Volume and Snapshot

Problem Statement:

As a developer, your development team has asked you to create a small Linux-based instance using a custom AMI for testing a new application.

Perform the following:

- Open the Amazon EC2 console
- Create a custom AMI
- Create the instance using a custom AMI
- Create an EBS Volume
- Attach an EBS Volume to a Linux Instance
- Create a Snapshot

